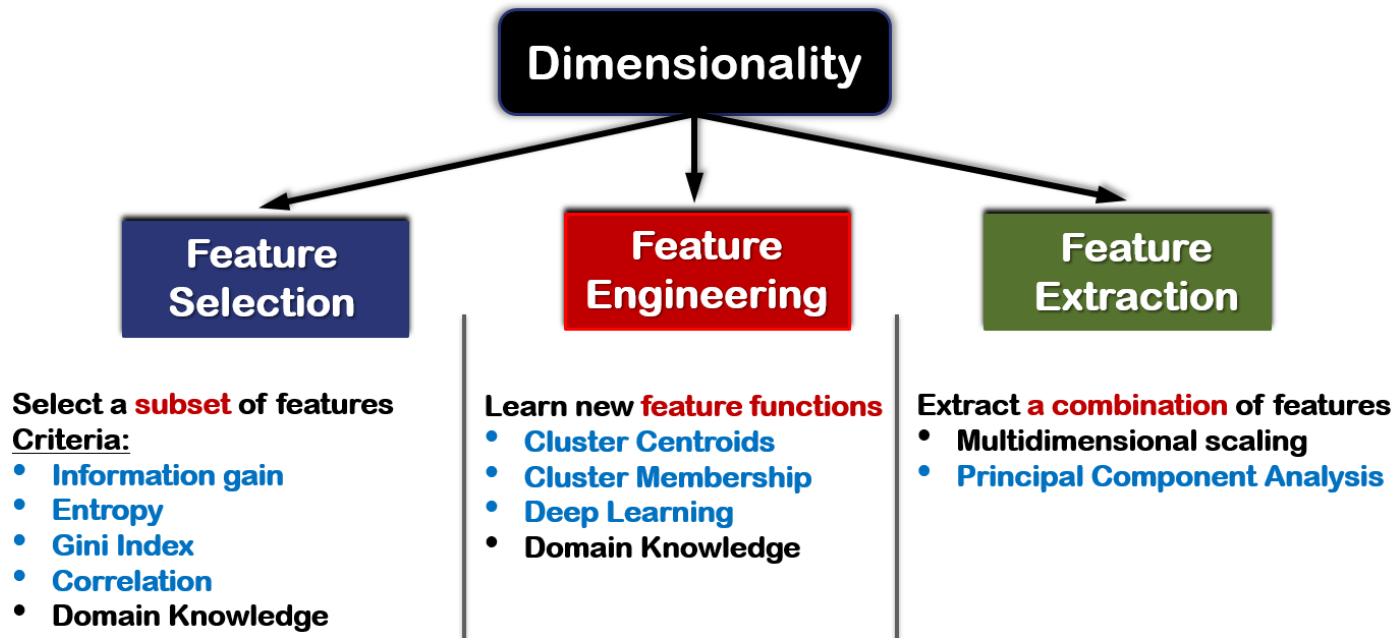


## PCA: Dimension Reduction



- Dimensionality and underdetermined problems
- Feature selection vs. feature extraction vs. feature engineering
- Dimension reduction (DR): unsupervised vs. supervised; linear vs. non-linear; orthogonal vs. non-orthogonal
- DR: linear, orthogonal with Principal Component Analysis (PCA)
- Eigenvalue and eigenvectors


**Prof. Nagiza F. Samatova**

samatova@csc.ncsu.edu

Department of Computer Science  
North Carolina State University

# Dimension Reduction

## MOTIVATION





# What Analysis Methods to Use?

Analysis methods fail for a few **gigabytes**.

## Method Complexity:

Calculate means  $O(n)$

Calculate Histogram  $O(n \log(n))$

Calculate PCA  $O(n \cdot d)$

Clustering algorithms  $O(n^2)$

If  $n = 10\text{GB}$ , then what is  $O(n)$  or  $O(n^2)$  on a teraflop computers?

$1\text{GB} = 10^9$  bytes  $1\text{Tflop} = 10^{12}$  op/sec

$O(\dots)$  - in the order of  
 $n$  – number of rows  
 $d$  - number of columns

| Data size $n$ | Algorithm Complexity |                 |                |
|---------------|----------------------|-----------------|----------------|
|               | $n$                  | $n \log(n)$     | $n^2$          |
| 100B          | $10^{-10}$ sec.      | $10^{-10}$ sec. | $10^{-8}$ sec. |
| 10KB          | $10^{-8}$ sec.       | $10^{-8}$ sec.  | $10^{-4}$ sec. |
| 1MB           | $10^{-6}$ sec.       | $10^{-5}$ sec.  | 1 sec.         |
| 100MB         | $10^{-4}$ sec.       | $10^{-3}$ sec.  | 3 hrs          |
| <b>10GB</b>   | $10^{-2}$ sec.       | 0.1 sec.        | <b>3 yrs.</b>  |


For illustration chart assumes  $10^{-12}$  sec.  
(**1Tflop/sec**) calculation time per data point

# How to Make Sense of Data?

*Know Your Limits & Be Smart!*



**Not humanly possible to browse a petabyte of data.**  
Analysis must reduce data to quantities of interest.



## Computations:

Must be smart about which probe combinations to see!


## Physical Experiments:

Must be smart about probe placement!

**To see 1 percent of a petabyte at 10 megabytes per second takes:**

**35 8-hour days!**

# It is not just the Size



# High-dimensional Data

- **Text Data**

- Record/Row: Document ID
- Dimension/Column: Each word in a collection of text documents

- **Image Data**

- Record/Row: Image ID
- Dimension/Column: Each pixel

- **Audio Data**

- Record/Row: Audio record ID
- Dimension/Column: Frequency of an audio signal OR a word if audio → text conversion

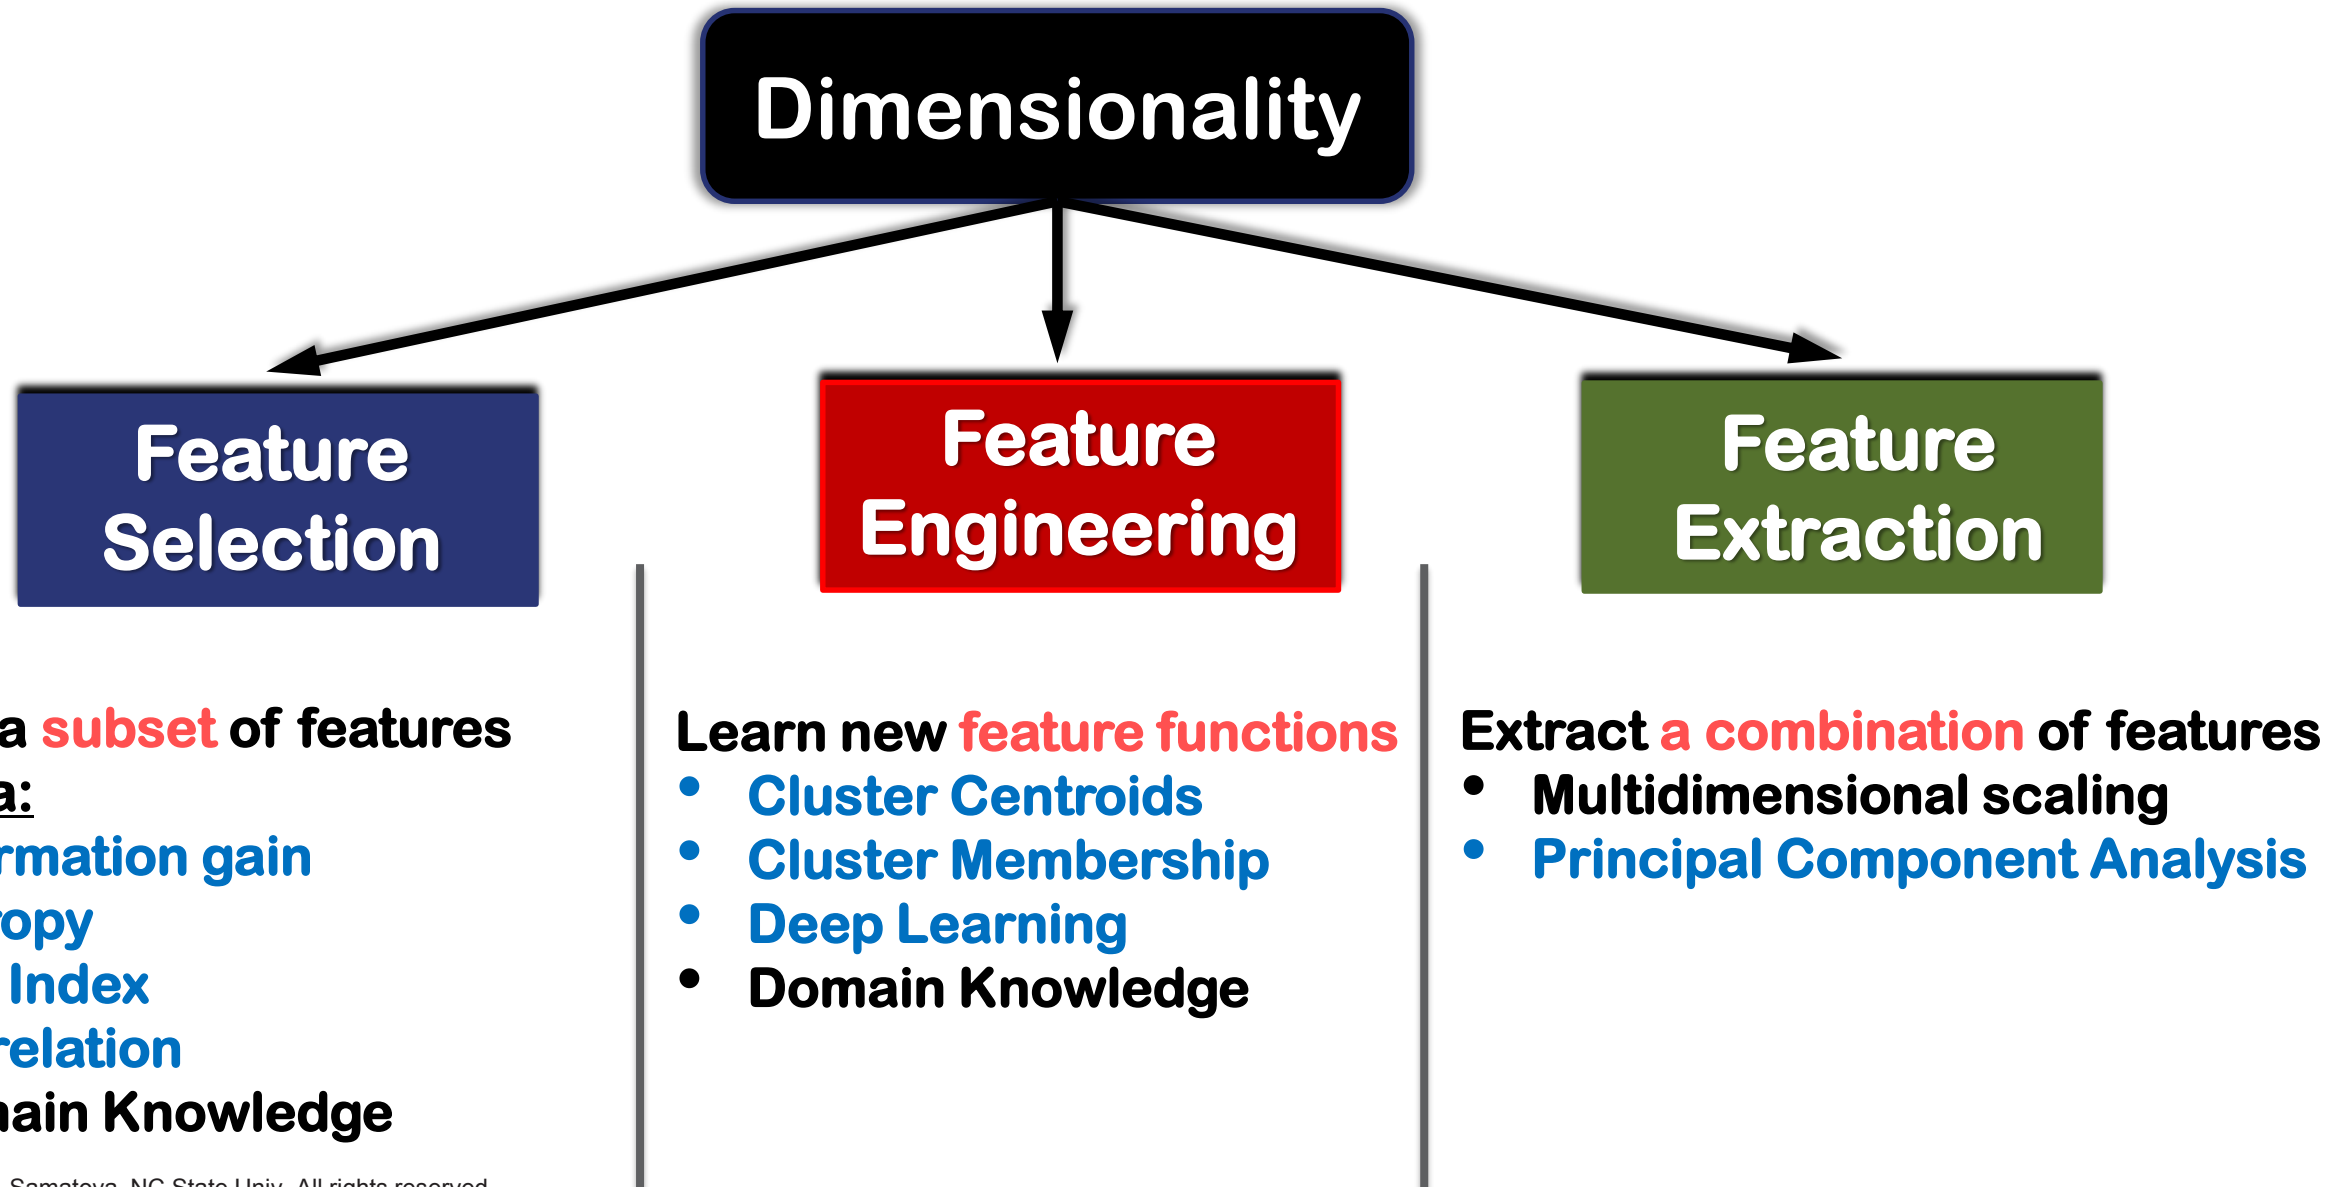
- **Sales Data**

- Record/Row: Transaction ID
- Dimension/Column: Each Product ID

# Taxonomy: Dimensionality-aware Methods

**FEATURE SELECTION, FEATURE  
EXTRACTION, FEATURE ENGINEERING**

# Dimensionality Challenge: Strategies to Cope With





# Motivation: Why to Cope with Dimensionality Issue?

- Not all the measured variables are important for understanding the underlying “interesting” phenomena:
  - could complicate the process of data analysis
- Decrease the **computational cost** for other data mining tasks:
  - Proximity measure calculations:  $O(d) \rightarrow O(k)$ ,  $d \gg k$
- Reduce the **noise** in the data:
  - improve signal to noise ratio
- Improve the **accuracy** of predictive models
  - better designed/engineered features capable of capturing non-linear signals in the data
- Reduce **collinearity** among variables/features
  - Critical for regression models

# What is Dimensionality Reduction (DR)?

**Objective:** To transform data from a **high-dimensional** representation to a **low-dimensional** representation, while **“best” preserving the information**.

Given dataset with  $n$  objects:

$$\begin{array}{ccc} X \in \mathbb{R}^{n \times m} & \xrightarrow{\text{DR}} & X' \in \mathbb{R}^{n \times p} \\ m \text{ dimensions} & & p \text{ dimensions} \end{array}$$


where  $p \ll m$

**Text Mining Example:**

- $m = 50,000$  words  $\rightarrow p \approx 100$  extracted features


# Example 1: $p=2 \rightarrow d=1$

Original data,  $p = 2$




$$X_{4 \times 2}$$

$x_1$ -projection,  $k=1$



$$X'_{4 \times 2} = X_{4 \times 2} \cdot \begin{bmatrix} 1 & 0 \\ 0 & 0 \end{bmatrix}$$

$x_2$ -projection,  $k = 1$



$$X'_{4 \times 2} = X_{4 \times 2} \cdot \begin{bmatrix} 0 & 0 \\ 0 & 1 \end{bmatrix}$$


$$X'_{m \times d} = X_{m \times d} \cdot P_{d \times d}$$

$$P_{d \times d} = \begin{bmatrix} 1 & 0 & \dots & 0 \\ 0 & 0 & \dots & 0 \\ \dots & \dots & \dots & \dots \\ 0 & 0 & 0 & 1 \end{bmatrix} - \text{projection matrix; some diagonal elements are 0}$$


Which projection is “**better**”?

## Example 2: Linear, Orthogonal Projection


Original data,  $d=2$



$x_1$ -projection,  $k=1$




$x_2$ -projection,  $k=1$




Is there a “better” projection?

Another basis,  $d=2$   
(rotate coord. system)



$e_1$ -projection,  
 $k=1$





Projection:

- Linear,  $e_1$  – line
- Orthogonal

$$e_1 \perp e_2$$

$e_1, e_2$  – eigenvectors of  $X$


# Example 3: Non-linear projection



**NON-linear** projection

# Example 4: Projection for Labeled Data

What is a “better” projection?



Ideal for machine learning (ML) algorithm:

- Points within the same group come from a Gaussian distribution (spherical shapes)
- Points from different groups are linearly separable

“better” means:

- it bests, ideally, **linearly separates** different groups of data, i.e.
- points from **the same group** are **closer** to each other and are **farther away** from the points in **different groups**

# Summary: Taxonomy of Dimension Reduction

- **Linear vs. NON-linear**
- **Orthogonal vs. non-orthogonal**
- **Unsupervised (unlabeled data) vs. supervised (labeled data)**

## Linear dimension reduction:

- Interpretable in original space
- Preserves non-linearity for visualization
- Orthogonal projections
- Non-orthogonal projections
- Favored for structure discovery

## Non-linear dimension reduction:

- Lower-dimensional representation
- Interpretable w.r.t. Non-linear transformation
- 1 to 3 orders of magnitude more computation
- Favored for prediction or classification

- **A lower-dimensional representation that contains the essence of the high dimensional data**
- **Blessing of dependence/correlation that saves us from the curse of dimensionality**

# Linear Orthogonal Dimension Reduction

**PCA: PRINCIPAL COMPONENT ANALYSIS**



# Dimensionality Challenge: Strategies to Cope With

## Dimensionality

```
graph TD; Dimensionality[Dimensionality] --> FeatureSelection[Feature Selection]; Dimensionality --> FeatureEngineering[Feature Engineering]; Dimensionality --> FeatureExtraction[Feature Extraction];
```

### Feature Selection

Select a **subset** of features

Criteria:

- Information gain
- Entropy
- Gini Index
- Correlation
- Domain Knowledge

### Feature Engineering

Learn new **feature functions**

- Cluster Centroids
- Cluster Membership
- Deep Learning
- Domain Knowledge

### Feature Extraction

Extract a **combination** of features

- Multidimensional scaling
- Principal Component Analysis

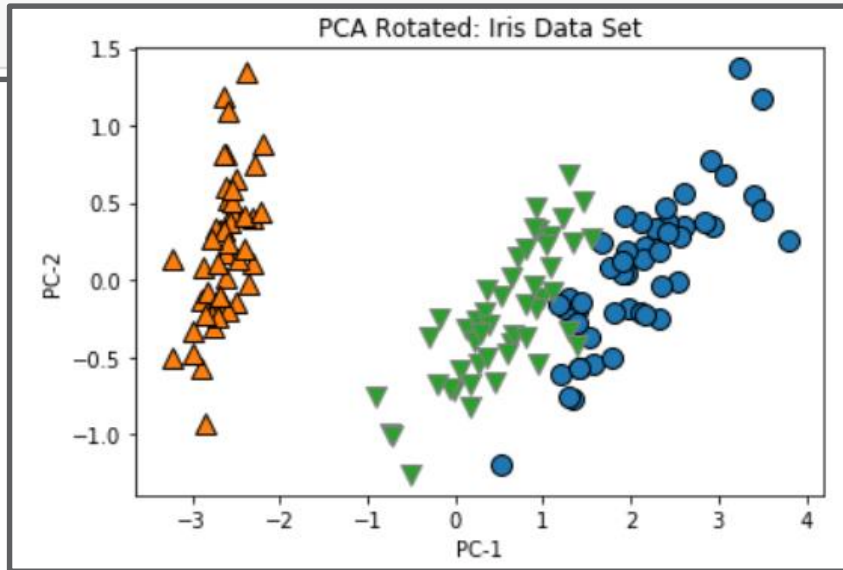
# What is “better” projection: $d \rightarrow k$ ( $k < d$ )?

- Many definitions are possible
- Definition 1:
  - Projection that **maximizes the VARIANCE** of the original  $d$ -dimensional data upon its projection onto the target  $k$ -dimensions ( $k < d$ )

# Python Code Example: Iris Data

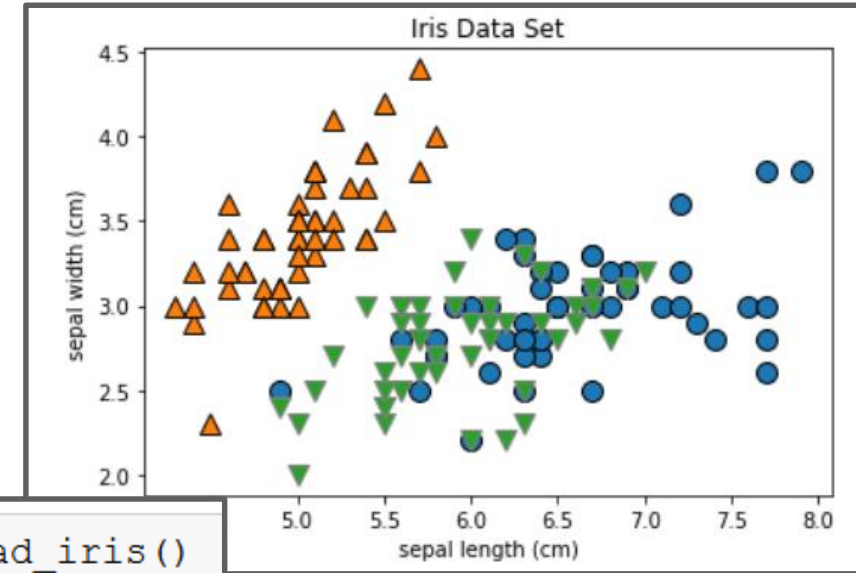
```
pca = decomposition.PCA(n_components=3)
pca.fit(X)
X_rot = pca.transform(X)
```

```
# Plot rotated 4-D iris data in 2D: the first two PCs
mglearn.discrete_scatter(X_rot[:, 0], X_rot[:, 1], y)
plt.title("PCA Rotated: Iris Data Set")
plt.xlabel("PC-1")
plt.ylabel("PC-2")
plt.show()
```



**PCA-transformed data in 2-d**

**original 4-d data in 2-d**

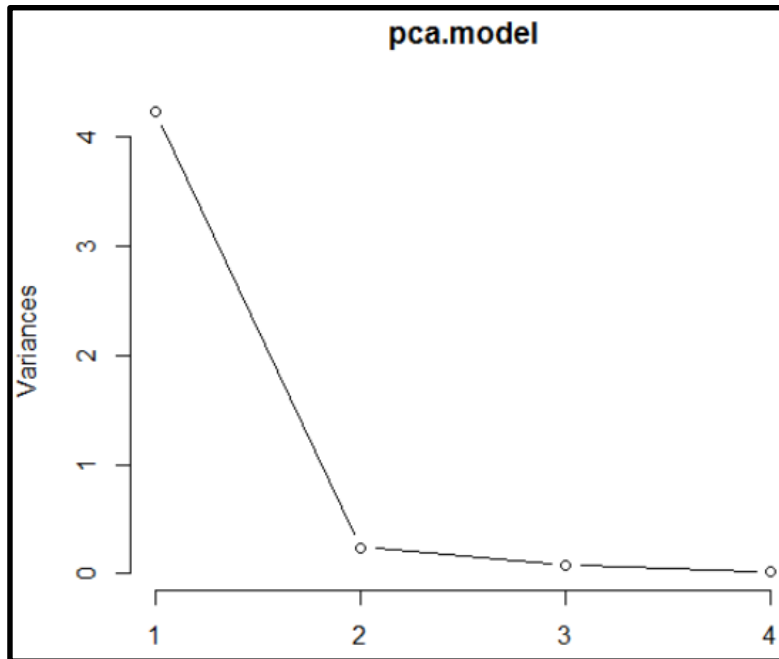


```
iris = datasets.load_iris()
X = iris.data
y = iris.target
print(iris.feature_names)
print(X[0:5])
print(iris.target_names)
print(y[[0, 51, 101]])
```

```
# Plot original 4-D iris data in 2D
mglearn.discrete_scatter(X[:, 0], X[:, 1], y)
plt.title("Iris Data Set")
plt.xlabel("sepal length (cm)")
plt.ylabel("sepal width (cm)")
plt.show()
```

# Visualizing PCA Results: Iris Data

## Screepplot



Successive variance accounted  
by each component

### Component loadings:

|              | Comp.1     | Comp.2     | Comp.3     | Comp.4     |
|--------------|------------|------------|------------|------------|
| Petal.Length | 0.5804131  | 0.02449161 | 0.1421264  | 0.8014492  |
| Petal.Width  | 0.5648565  | 0.06694199 | 0.6342727  | -0.5235971 |
| Sepal.Length | 0.5210659  | 0.37741762 | -0.7195664 | -0.2612863 |
| Sepal.Width  | -0.2693474 | 0.92329566 | 0.2443818  | 0.1235096  |

### Component variances:

|                        | Comp.1    | Comp.2    | Comp.3     | Comp.4      |
|------------------------|-----------|-----------|------------|-------------|
| Standard deviation     | 1.7083611 | 0.9560494 | 0.38308860 | 0.143926497 |
| Proportion of Variance | 0.7296245 | 0.2285076 | 0.03668922 | 0.005178709 |
| Cumulative Proportion  | 0.7296245 | 0.9581321 | 0.99482129 | 1.000000000 |

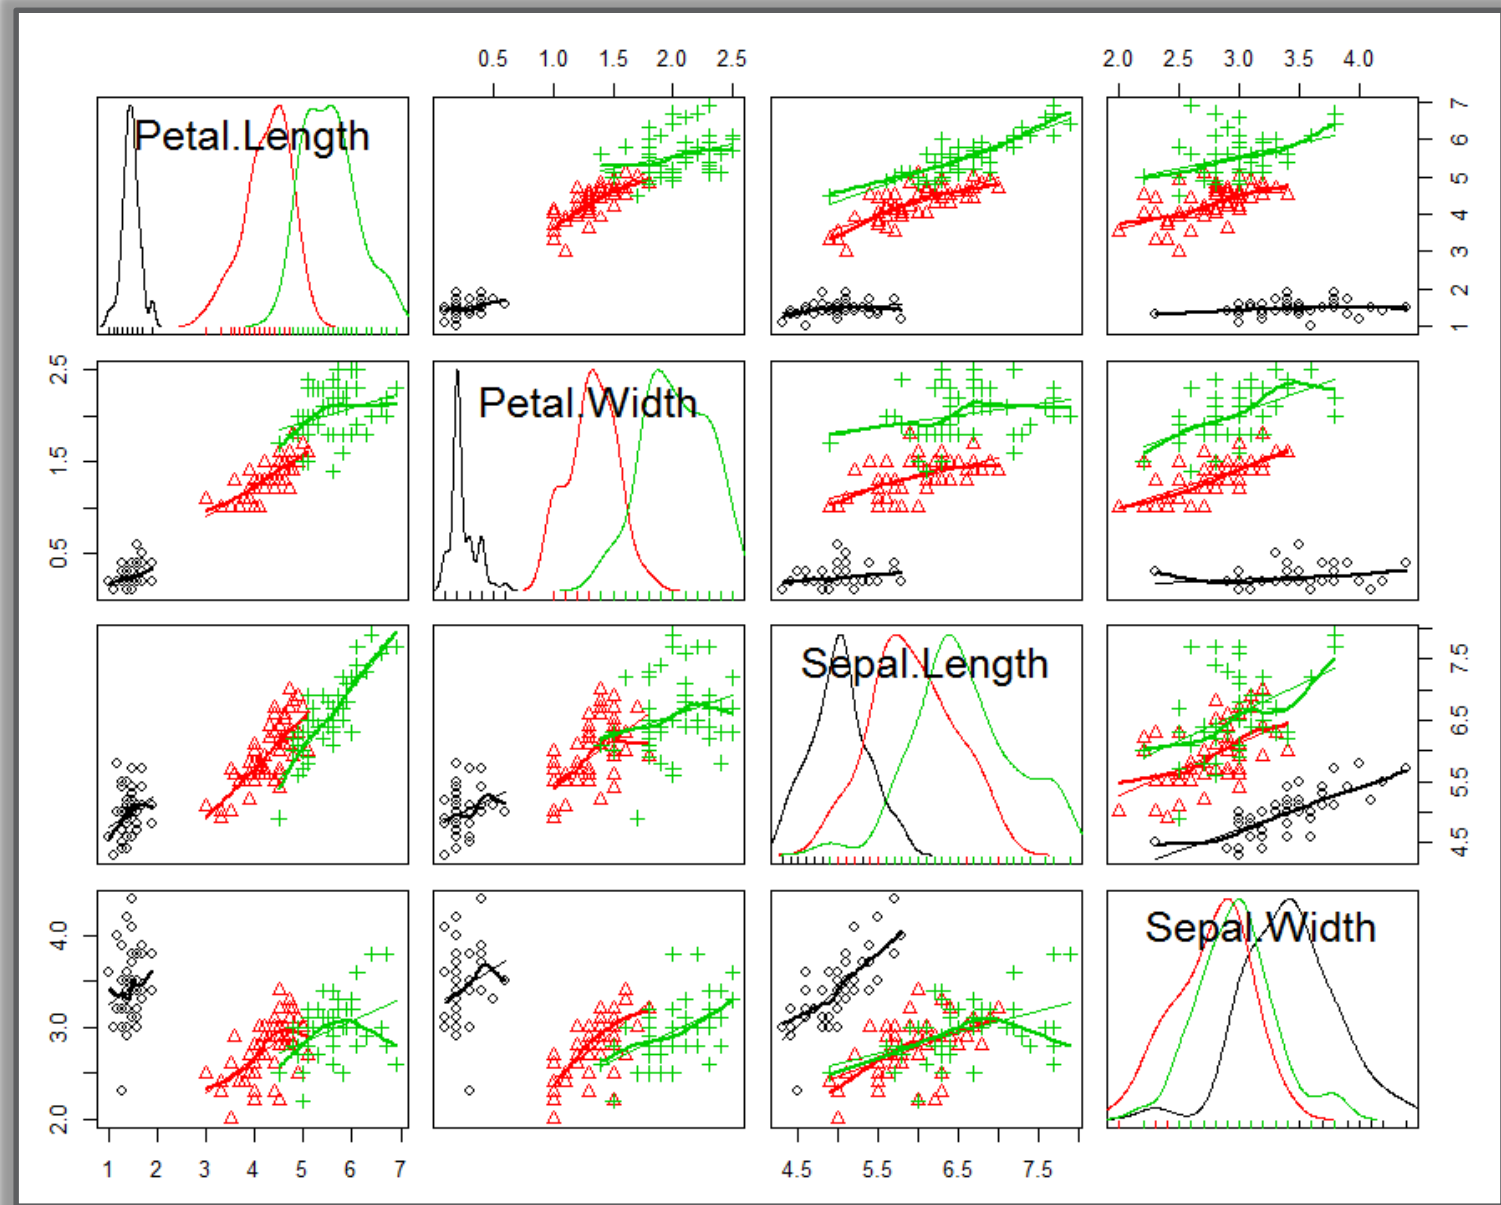
### Importance of components:

|                        | Comp.1    | Comp.2    | Comp.3     | Comp.4      |
|------------------------|-----------|-----------|------------|-------------|
| Standard deviation     | 1.7083611 | 0.9560494 | 0.38308860 | 0.143926497 |
| Proportion of Variance | 0.7296245 | 0.2285076 | 0.03668922 | 0.005178709 |
| Cumulative Proportion  | 0.7296245 | 0.9581321 | 0.99482129 | 1.000000000 |




# Original Relationships: Features are **Correlated!**

## Scatterplot Matrix



# Extracted Principal Components (PCs) are **Uncorrelated**

|     | PC1    | PC2    | PC3    | PC4    |
|-----|--------|--------|--------|--------|
| PC1 | 1e+00  | -1e-16 | -1e-15 | 2e-15  |
| PC2 | -1e-16 | 1e+00  | 2e-16  | -5e-16 |
| PC3 | -1e-15 | 2e-16  | 1e+00  | -1e-15 |
| PC4 | 2e-15  | -5e-16 | -1e-15 | 1e+00  |



The new extracted features (PCs) are **UNCORRELATED!**

# PC is a weighted linear sum of original features

Extracted Feature → PCA-based Feature Extraction:

$$PC = w_1 * f_1 + w_2 * f_2 + \dots + w_d * f_d$$

The magnitude of each **weight** indicates **how important the corresponding feature is** → it could be used as a **feature selection** technique!

|              | PC1   | PC2   | PC3   | PC4  |
|--------------|-------|-------|-------|------|
| Sepal.Length | 0.36  | -0.66 | 0.58  | 0.3  |
| Sepal.Width  | -0.08 | -0.73 | -0.60 | -0.3 |
| Petal.Length | 0.86  | 0.17  | -0.08 | -0.5 |
| Petal.Width  | 0.36  | 0.08  | -0.55 | 0.8  |

$$PC1 = 0.86 \text{ Petal.Length} + 0.36 \text{ Petal.Width} \\ + 0.36 \text{ Sepal.Length} - 0.08 \text{ Sepal.Width}$$

$$PC2 = -0.73 \text{ Sepal.Width} - 0.66 \text{ Sepal.Length} \\ + 0.17 \text{ Petal.Length} + 0.08 \text{ Petal.Width}$$

# Preserved Variability for **Top-k PCs**: $d \rightarrow k$ ( $k \ll d$ )

**Percentage of variability** preserved  
if the first  $k$  PCs are used for projection:

$$\frac{\sum_{i=1}^k \lambda_i}{\sum_{j=1}^d \lambda_j}$$

Standard deviation = sqrt (Variance)


```
Standard deviations (1, ..., p=4):  
[1] 2.06 0.49 0.28 0.15
```

**Variance preserved by PC1:**

$$2.06 \times 2.06 = 4.2$$

**SORTED Eigenvalues:** `pca.model$sdev`

$$\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_k \geq \dots \geq \lambda_d \geq 0$$





# Eigenvalues: Importance of Eigenvectors

**SORTED:**

$$\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_k \geq \dots \geq \lambda_d \geq 0$$

Importance of Principal Components (PC), or Eigenvectors, or Rotations

**PC1 preserves more variance than PC2**

PC2 preserves more variance than PC3

....




**Proportion of Variance Preserved** if only  $k$  PCs are used:

$$\frac{\sum_{i=1}^k \lambda_i}{\sum_{j=1}^d \lambda_j}$$

# Principal Components (eigenvectors) are NOT unique

PCA: **Linear Orthogonal** Transformation:

- PCs are simply a rotation of the original coordinate system
- Each PC is a **line**
- Each PC is **perpendicular** to the other PCs
  - PCs are **Uncorrelated**
  - The angle between any pair of PCs is  $90^\circ$




# PCA: Linear Orthogonal Dimension Reduction

Principal Component Analysis (**PCA**) finds **intrinsic** dimensionality and allows for low-dimensional representation of the data.



Principal Component Analysis is the **Spectral Decomposition of the Covariance Matrix**.



# PCA: Covariance vs. Correlation Matrix

- PCA results **depend on the scales** at which the variables are measured:
  - PCA should only be used with the raw data if all variables have the same units of measure
- PCA results **depend on the variances** of the variables: the ones with the highest sample variances will tend to be emphasized in the first few principal components:
  - Use PCA with covariance matrix only if you wish to give variables with higher variances more weight in the analysis
- If the variables either have **different units of measurement** (i.e., pounds, feet, gallons, etc), or if we wish **each variable to receive equal weight in the analysis**, then the variables should be **standardized (Z-scores)** before a principal components analysis is carried out.