

Matrix Transformations & Decomposition

Centering and standardizing, correlation and covariance matrix, projection and rotation matrices, matrix decomposition, Principal Component Analysis (PCA), Singular Value Decomposition (SVD), eigenvalues, eigenvectors

Nagiza F. Samatova, samatova@csc.ncsu.edu

Professor, Department of Computer Science
North Carolina State University

Vector Transformation

CENTERING AND STANDARDIZING

Centering and Standardizing

Let $x = (x_1, x_2, \dots, x_n)$ be the column vector

The **mean** of the vector:

$$\bar{x} = \frac{x_1 + x_2 + \dots + x_n}{n} = \frac{1}{n} \sum_{i=1}^n x_i$$

$$\bar{x} = \frac{2.1 + 2.5 + 4.0 + 3.6}{4} = 3.05$$

Example in R:

```
> x <- c(2.1, 2.5, 4.0, 3.6)
> mean(x)
[1] 3.05
```

Example in Python:

```
1 x = np.array([2.1, 2.5, 4.0, 3.6])
2 np.mean(x)

3.05
```

Economic Growth % (x_i)	S & P 500 Returns % (y_i)
2.1	8
2.5	12
4.0	14
3.6	10

Centering

Let $x = (x_1, x_2, \dots, x_n)$ be the column vector

The **mean** of the vector: $\bar{x} = \frac{x_1 + x_2 + \dots + x_n}{n} = \frac{1}{n} \sum_{i=1}^n x_i$ ← scalar, number

Centering the vector:
center x at its mean

$$x_c = x - \bar{x} = (x_1 - \bar{x}, x_2 - \bar{x}, \dots, x_n - \bar{x}) \quad \leftarrow \text{vector}$$

Note: The mean of the centered vector is zero: $\overline{x_c} = 0$

$$\bar{x} = \frac{2.1 + 2.5 + 4.0 + 3.6}{4} = 3.05$$

$$x_c = (2.1 - 3.05, 2.5 - 3.05, 4.0 - 3.05, 3.6 - 3.05) \\ = (-.95, -0.55, 0.95, 0.55)$$

Economic Growth % (x_i)	S & P 500 Returns % (y_i)
2.1	8
2.5	12
4.0	14
3.6	10

Centering in Python and in R

Centering the vector:
center x at its mean

$$x_c = x - \bar{x} = (x_1 - \bar{x}, x_2 - \bar{x}, \dots, x_n - \bar{x})$$

← vector

Example in R:

```
> x <- c(2.1, 2.5, 4.0, 3.6)
> mean(x)
[1] 3.05
> xc <- x - mean(x)
> xc
[1] -0.95 -0.55  0.95  0.55
```

Centering the vector:
center x at its mean

Example in Python:

```
1 xc = x - np.mean(x)
2 xc
array([-0.95, -0.55,  0.95,  0.55])
```

Economic Growth % (x_i)	S & P 500 Returns % (y_i)
2.1	8
2.5	12
4.0	14
3.6	10

Standardizing in R

Let $x = (x_1, x_2, \dots, x_n)$ be the column vector

Centered vector:

$$x_c = x - \bar{x} = (x_1 - \bar{x}, x_2 - \bar{x}, \dots, x_n - \bar{x})$$

Variance: $var(x) = \frac{(x_1 - \bar{x})^2 + (x_2 - \bar{x})^2 + \dots + (x_n - \bar{x})^2}{n-1}$

Standard Deviation: $sd(x) = \sqrt{var(x)}$

Standardizing using standard deviation:

$$x_s = \frac{x}{sd(x)}$$

Standardizing using mean & standard deviation (Z-score):

$$\text{Z-score} = \frac{x - \bar{x}}{sd(x)} = \frac{x_c}{sd(x)}$$

Economic Growth % (x_i)	S & P 500 Returns % (y_i)
2.1	8
2.5	12
4.0	14
3.6	10

Example in R:

```
> x <- c(2.1, 2.5, 4.0, 3.6)
> var(x)
[1] 0.80333
> sd(x)
[1] 0.89629
> sd(x) * sd(x)
[1] 0.80333
> xc * xc
[1] 0.9025 0.3025 0.9025 0.3025
> sum(xc * xc)
[1] 2.41
> sum(xc * xc) / (4-1)
[1] 0.80333
> xs <- x / sd(x)
> xs
[1] 2.3430 2.7893 4.4628 4.0166
> Z.score <- xc / sd(x)
> Z.score
[1] -1.05993 -0.61364 1.05993 0.61364
```

Standardizing in Python

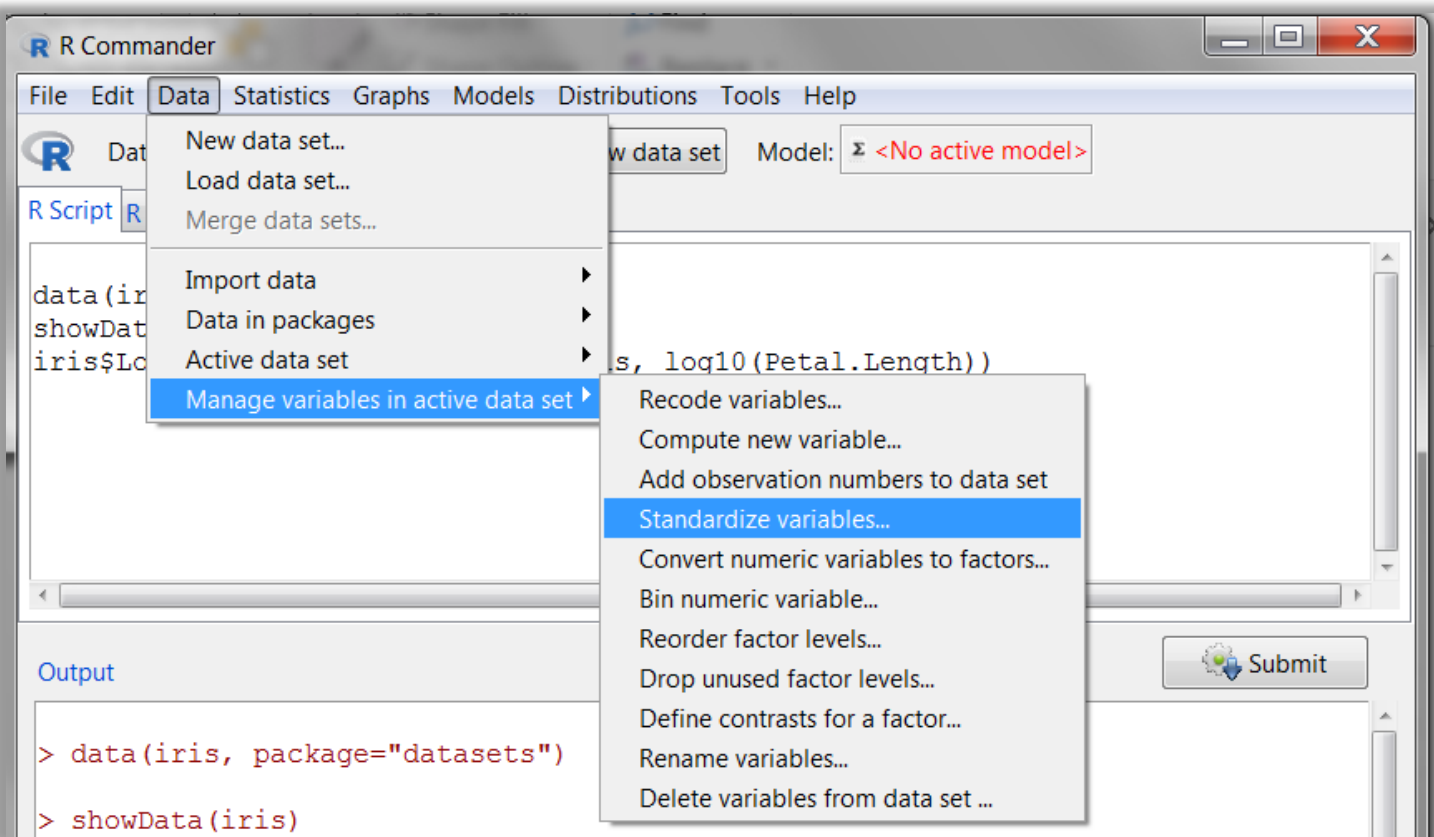
```
1 x = np.array([2.1,2.5,4.0,3.6])
2 var_x = np.var(x,ddof=1)
3 sd = np.std(x,ddof=1)
4 sd_squared = sd**2
5 xc_squared = xc**2
6 sum_xc_squared = np.sum(xc_squared)
7 var = sum_xc_squared/(len(x)-1)
8 xs = x/sd
9 z_score = xc/sd
10
11 print(round(var_x, 5))
12 print(round(sd, 5))
13 print(round(sd_squared, 5))
14 print(xc_squared)
15 print(round(sum_xc_squared, 2))
16 print(round(var, 5))
17 print(xs)
18 print(z_score)
```

Economic Growth % (x_i)	S & P 500 Returns % (y_i)
2.1	8
2.5	12
4.0	14
3.6	10

Output:

```
0.80333
0.89629
0.80333
[0.9025 0.3025 0.9025 0.3025]
2.41
0.80333
[2.34299521 2.78928001 4.46284802 4.01656322]
[-1.0599264 -0.6136416 1.0599264 0.6136416]
```

Apply Transformations: Standardize Variables



The screenshot shows a window titled 'iris' displaying a data frame with 7 columns: Species, Log10.Petal.Length, Z.Petal.Length, Z.Petal.Width, Z.Sepal.Length, Z.Sepal.Width, and Z.Petal.Length. The first five rows are shown, all belonging to the 'setosa' species.

	Species	Log10.Petal.Length	Z.Petal.Length	Z.Petal.Width	Z.Sepal.Length	Z.Sepal.Width
1	setosa	0.14612804	-1.33575163	-1.3110521482	-0.89767388	1.0
2	setosa	0.14612804	-1.33575163	-1.3110521482	-1.13920048	-0.1
3	setosa	0.11394335	-1.39239929	-1.3110521482	-1.38072709	0.3
4	setosa	0.17609126	-1.27910398	-1.3110521482	-1.50149039	0.0
5	setosa	0.14612804	-1.33575163	-1.3110521482	-1.01843718	1.2


Other Types of Matrices

CORRELATION AND COVARIANCE MATRIX

Covariance and Correlation

Covariance and correlation describe how two variables are related.

- Variables are positively related if they move in the same direction.
- Variables are inversely related if they move in opposite directions.



Covariance: Formula

$x = (x_1, x_2, \dots, x_n)$: the independent variable

$y = (y_1, y_2, \dots, y_n)$: the dependent variable

\bar{x} : the mean of x

\bar{y} : the mean of y

n : the number of points in the sample

$$\text{cov}(x, y) = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{n - 1}$$

$$\text{cov}(x, y) = \frac{1}{n - 1} x_c^T y_c$$

Cross-product (inner product) of centered variables normalized by the sample size minus 1 ($n - 1$).

Covariance: Example in Python

centering

centering

normalized
cross-product
covariance

Economic Growth % (x_i)	S & P 500 Returns % (y_i)
2.1	8
2.5	12
4.0	14
3.6	10

```
1 x = np.array([2.1,2.5,4.0,3.6])
2 xc = x - np.mean(x)
3 y = np.array([8,12,14,10])
4 yc = y - np.mean(y)
5
6 norm_cross_prod = np.sum(np.matmul(xc.T,yc))/(len(y)-1)
7 print(norm_cross_prod)
8 np.cov(x,y)[1,0]
```

```
1.5333333333333332
```

```
1.5333333333333332
```

Covariance: Example in R

$$\text{cov}(x, y) = \frac{1}{n-1} x_c^T y_c$$

Cross-product (inner product) of centered variables normalized by the sample size minus 1 ($n - 1$).

$$\text{cov}(x, y) = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{n-1}$$

Economic Growth % (x_i)	S & P 500 Returns % (y_i)
2.1	8
2.5	12
4.0	14
3.6	10

centering
centering
normalized
cross-product
covariance

```
> x <- c(2.1, 2.5, 4.0, 3.6)
> xc <- x - mean(x)
> y <- c(8, 12, 14, 10)
> yc <- y - mean(y)
>
> sum (t(xc) %*% yc) / (length(y) - 1)
[1] 1.5333
> cov (x, y)
[1] 1.5333
```


Correlation: Formula

$$cor(x, y) = \frac{cov(x, y)}{sd(x) \times sd(y)}$$

- **Correlation** is a scaled version of covariance (i.e., scaled by the standard deviations)
- Correlation and covariance always have the same sign (positive, negative, or 0)
 - when the sign is positive, the variables are said to be **positively correlated**
 - when the sign is negative, the variables are said to be **negatively correlated**
 - and when the sign is 0, the variables are said to be **uncorrelated**
- Correlation is dimensionless, since the numerator and denominator have the same physical units.
- Correlation will always take on a value between 1 and – 1:
 - If the correlation coefficient is **+1**, the variables have a perfect positive correlation. This means that if one variable moves a given amount, the second moves proportionally in the same direction.
 - If correlation coefficient is **–1**, the variables are perfectly negatively correlated (or inversely correlated) and move in opposition to each other. If one variable increases, the other variable decreases proportionally.
 - If correlation coefficient is **zero**, no relationship exists between the variables. If one variable moves, no predictions about the movement of the other variable can be made.

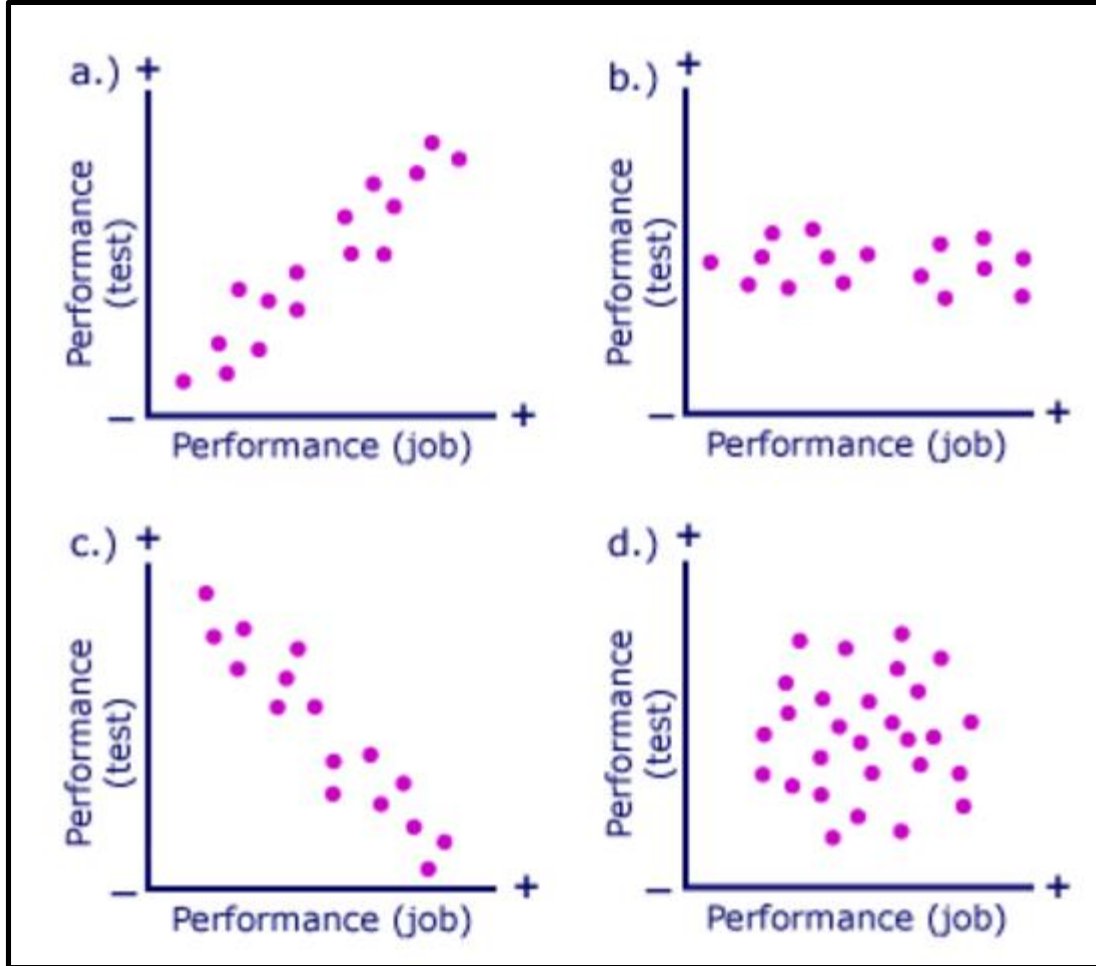
Understanding Correlation (correlation.R)

```
> x <- c(1:10)
> blue <- c(1:10)
> red <- c(3:12)
> purple <- sort(red, decreasing=TRUE)
>
> plot(x,blue,col="blue",pch=19)
> points(red,col="red",pch=19)
> points(purple,col="purple",pch=19)
>
> cor(blue,red)
[1] 1
> cor(red,purple)
[1] -1
```



```
> noise <- abs(rnorm(10, mean=5))
> noise
[1] 5.0548 6.2234 4.0714 3.4522 4.9278 4.6648 4.8472 5.1168 4.3224 5.5515
> points(noise, col="grey", pch=19)
> cor(blue,noise)
[1] 0.0073982
```

Correlation: Examples



In each of the graphs, are job performance and test performance shown to be positively related, inversely related, or unrelated?

Answers:

- a) positively related**
- b) unrelated**
- c) inversely related**
- d) unrelated**

Exercise: Compute Covariance & Correlation

Month	Return of Stock A	Return of Market Index
1	2.3	1.3
2	2.5	5.0
3	1.9	0.8
4	2.4	1.9
5	2.1	1.1

1. Using the table, show your calculations and R codes for computing the correlation of Stock A's returns and the return of the market index.

2. Do the same for the covariance.

Exercise: Solution

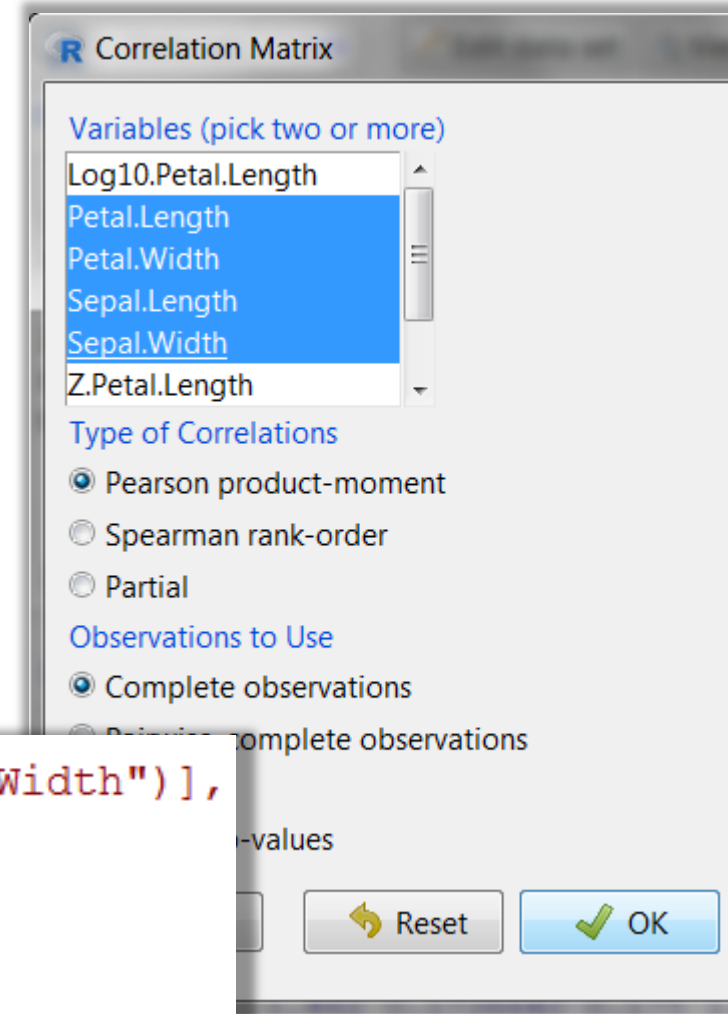
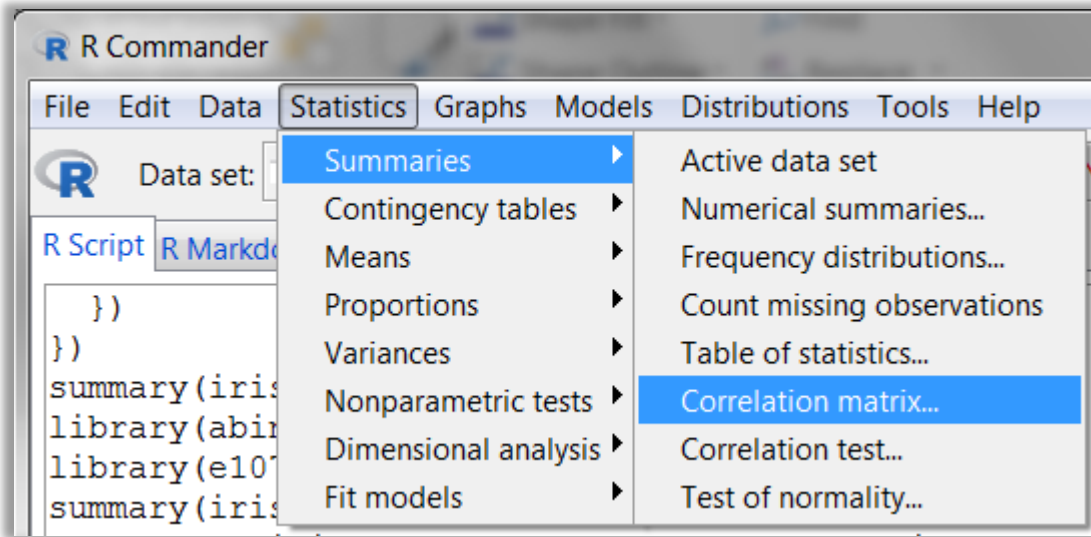
Month	Return of Stock A	Return of Market Index
1	2.3	1.3
2	2.5	5.0
3	1.9	0.8
4	2.4	1.9
5	2.1	1.1

			step 1	step 2
	Stock A	Market Return	$(x_i - \bar{x})^2$	$(y_i - \bar{y})^2$
	2.30	1.30	0.0036	0.5184
	2.50	5.00	0.0676	8.8804
	1.90	0.80	0.1156	1.4884
	2.40	1.90	0.0256	0.0144
	2.10	1.10	0.0196	0.8464
Sum			0.2320	11.7480
Average	2.24	2.02		
Sum ÷ 4			0.0580	2.9370
Standard deviation			0.2408	1.7138

$$\begin{aligned} \text{cor}(x, y) &= \frac{\text{cov}(x, y)}{\text{sd}(x)\text{sd}(y)} = \\ &= \frac{0.31}{(0.24)(1.71)} = \frac{0.31}{0.41} = 0.76 \end{aligned}$$

```
> x <- c(2.3,2.5,1.9,2.4,2.1)
> y <- c(1.3,5.0,0.8,1.9,1.1)
> mean(x)
[1] 2.24
> mean(y)
[1] 2.02
> sd(x)
[1] 0.24083
> sd(y)
[1] 1.7138
> cov(x,y)
[1] 0.314
> cor(x,y)
[1] 0.76079
> cov(x,y) / (sd(x)*sd(y))
[1] 0.76079
```

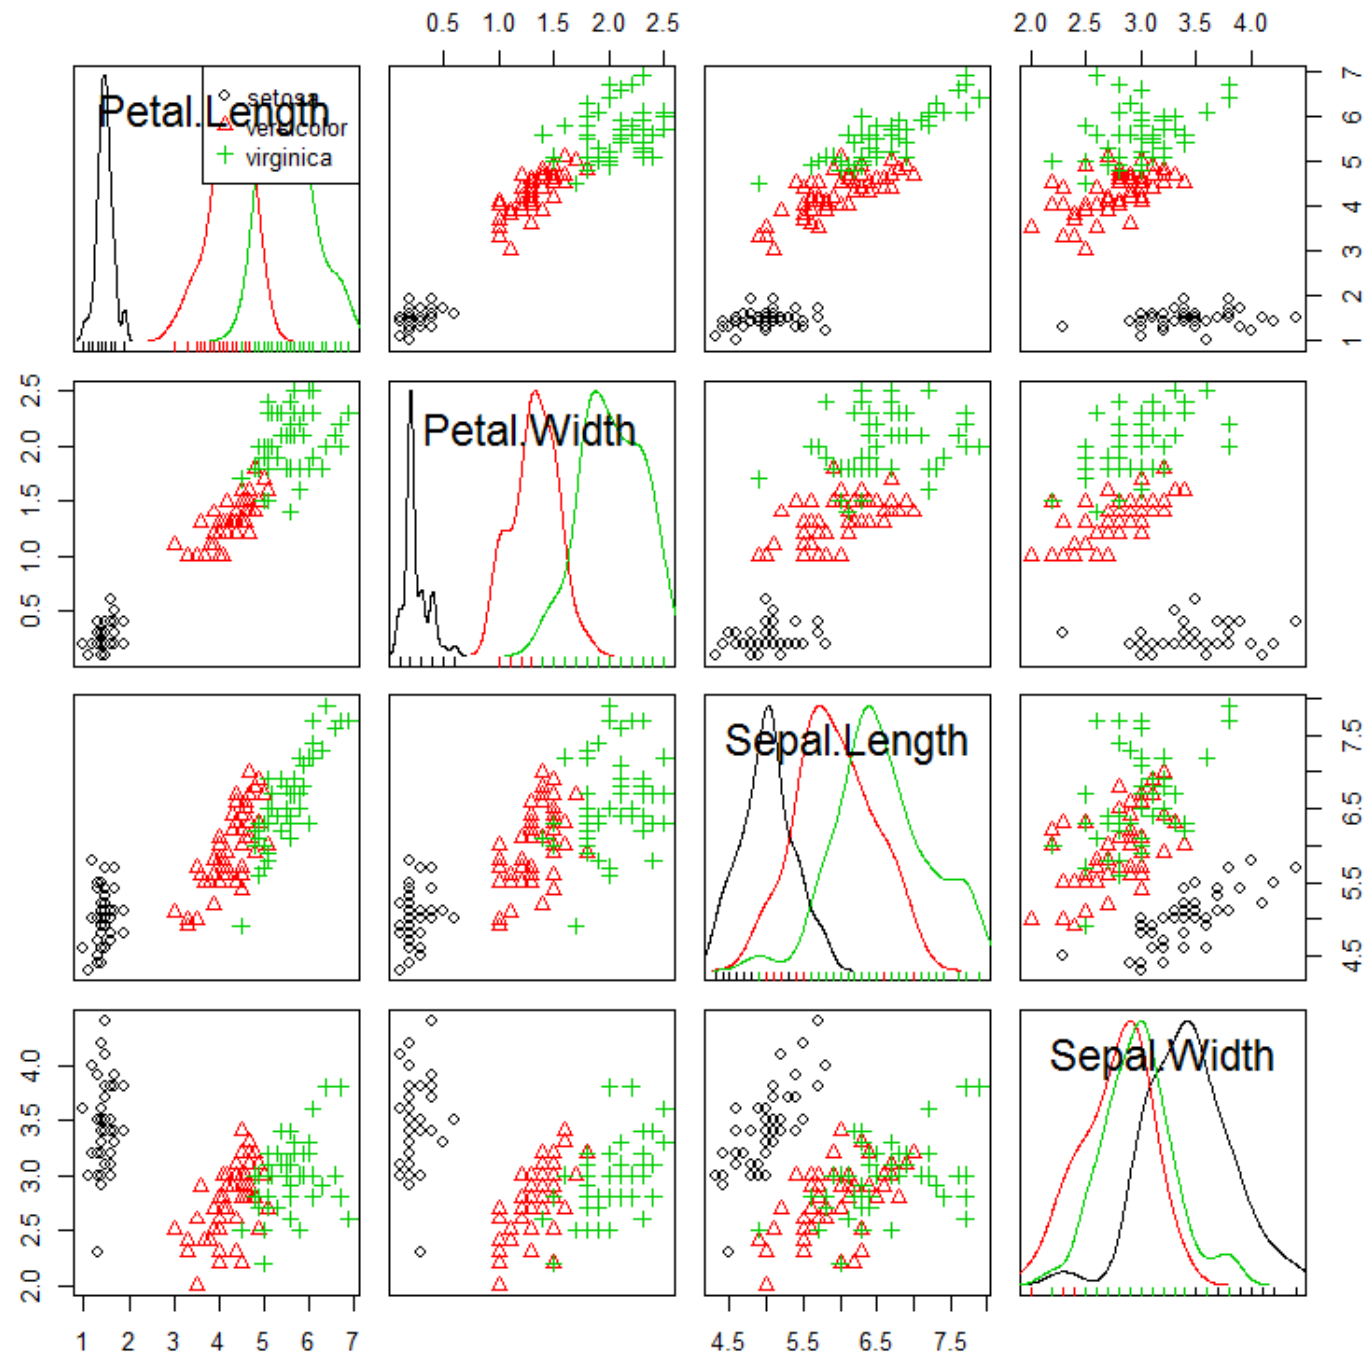
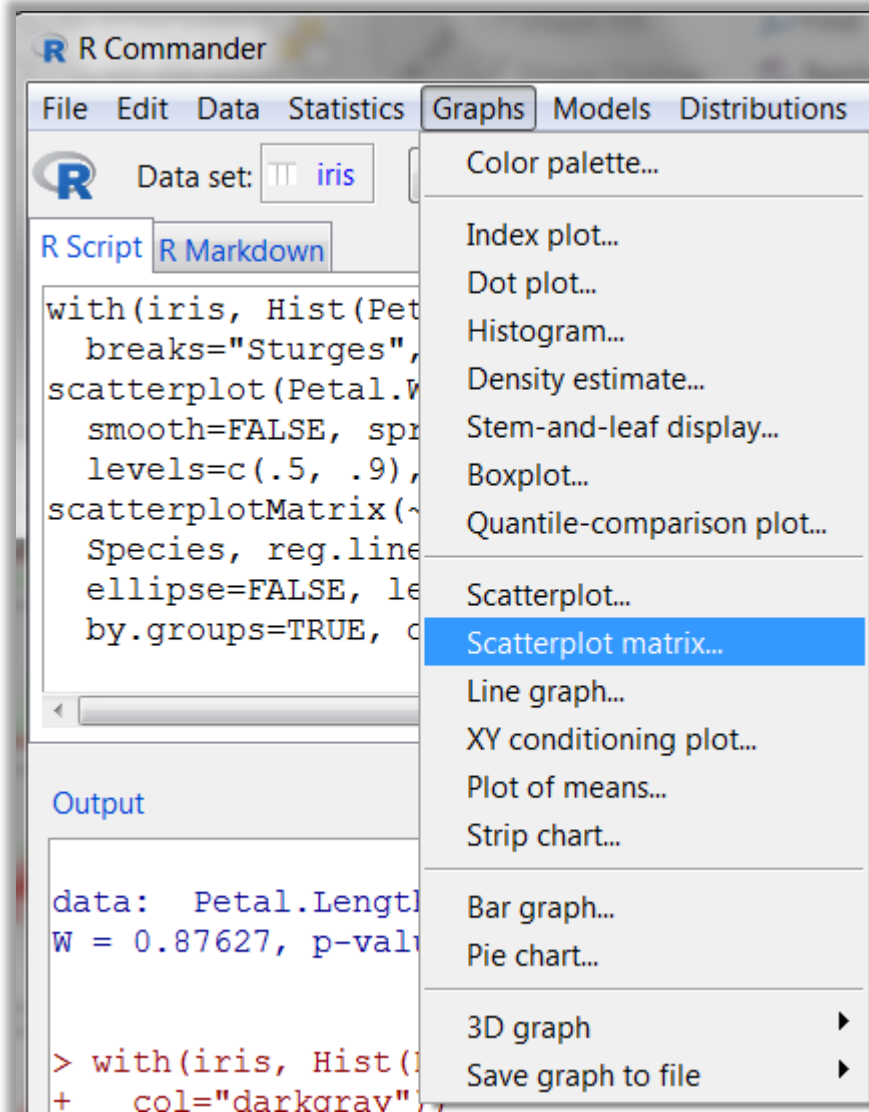
Correlation Matrix with R Commander



```
> cor(iris[,c("Petal.Length", "Petal.Width", "Sepal.Length", "Sepal.Width")],  
+     use="complete")
```

	Petal.Length	Petal.Width	Sepal.Length	Sepal.Width
Petal.Length	1.0000000	0.9628654	0.8717538	-0.4284401
Petal.Width	0.9628654	1.0000000	0.8179411	-0.3661259
Sepal.Length	0.8717538	0.8179411	1.0000000	-0.1175698
Sepal.Width	-0.4284401	-0.3661259	-0.1175698	1.0000000

Scatterplot Matrix

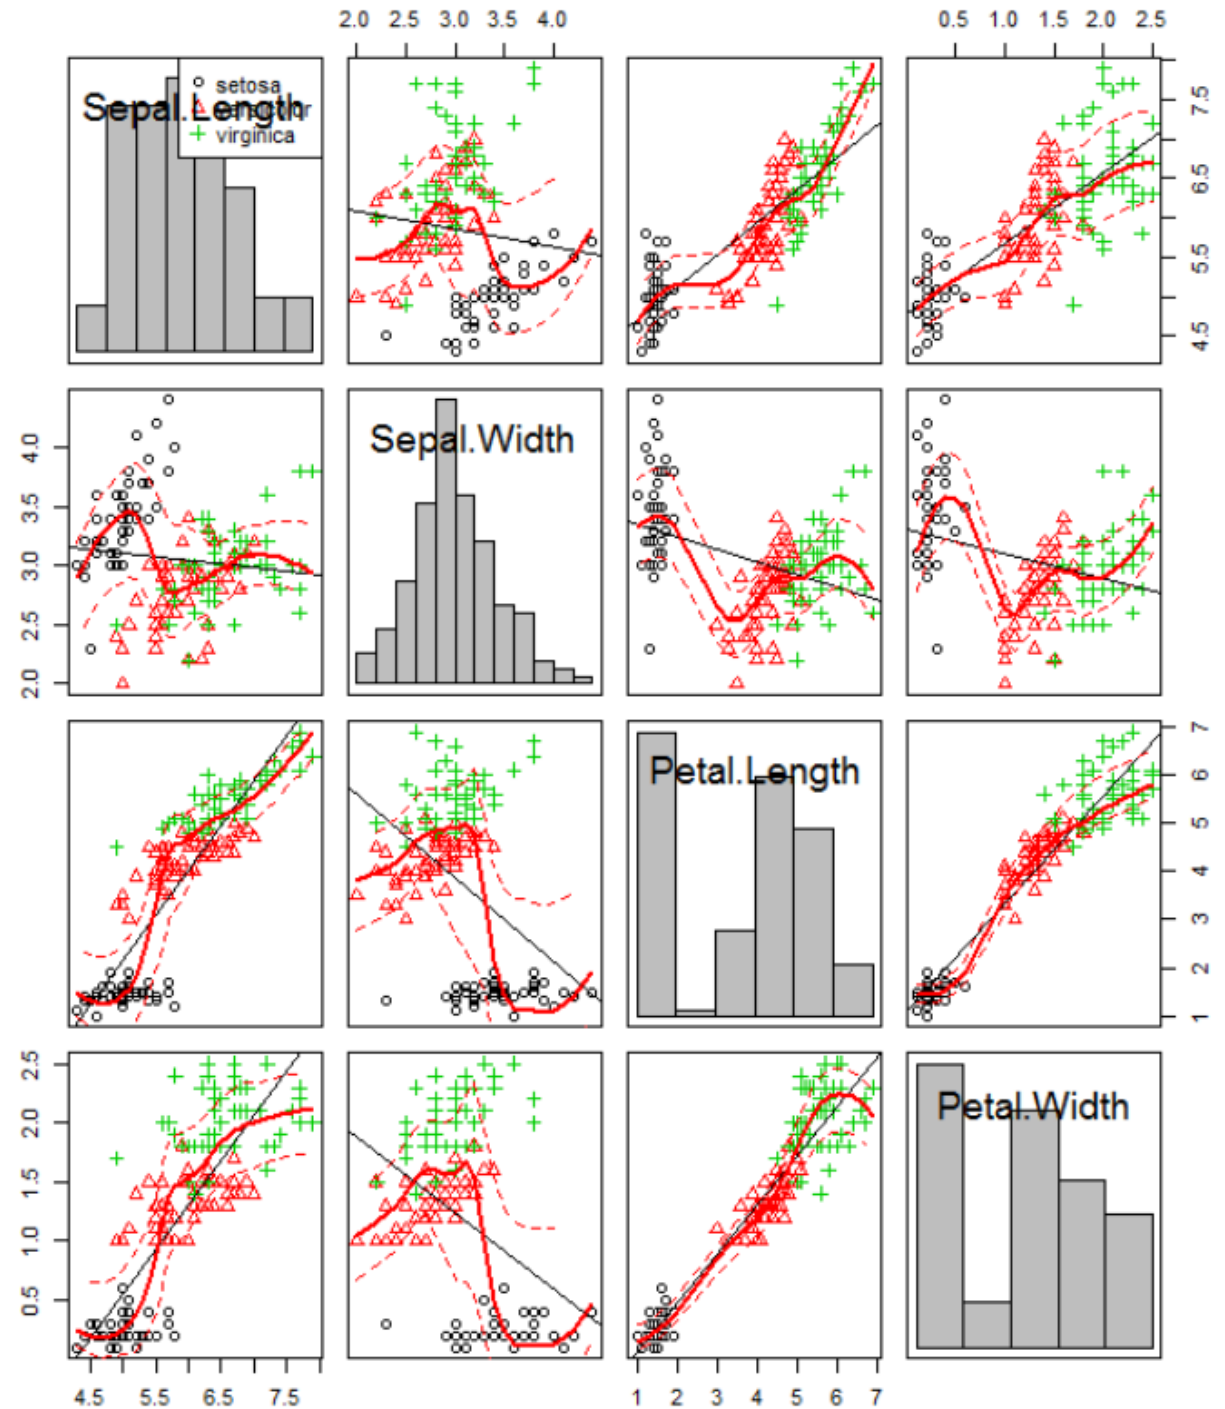


Visualizing Relationships: Continuous Variables

```
install.packages ("car")
```

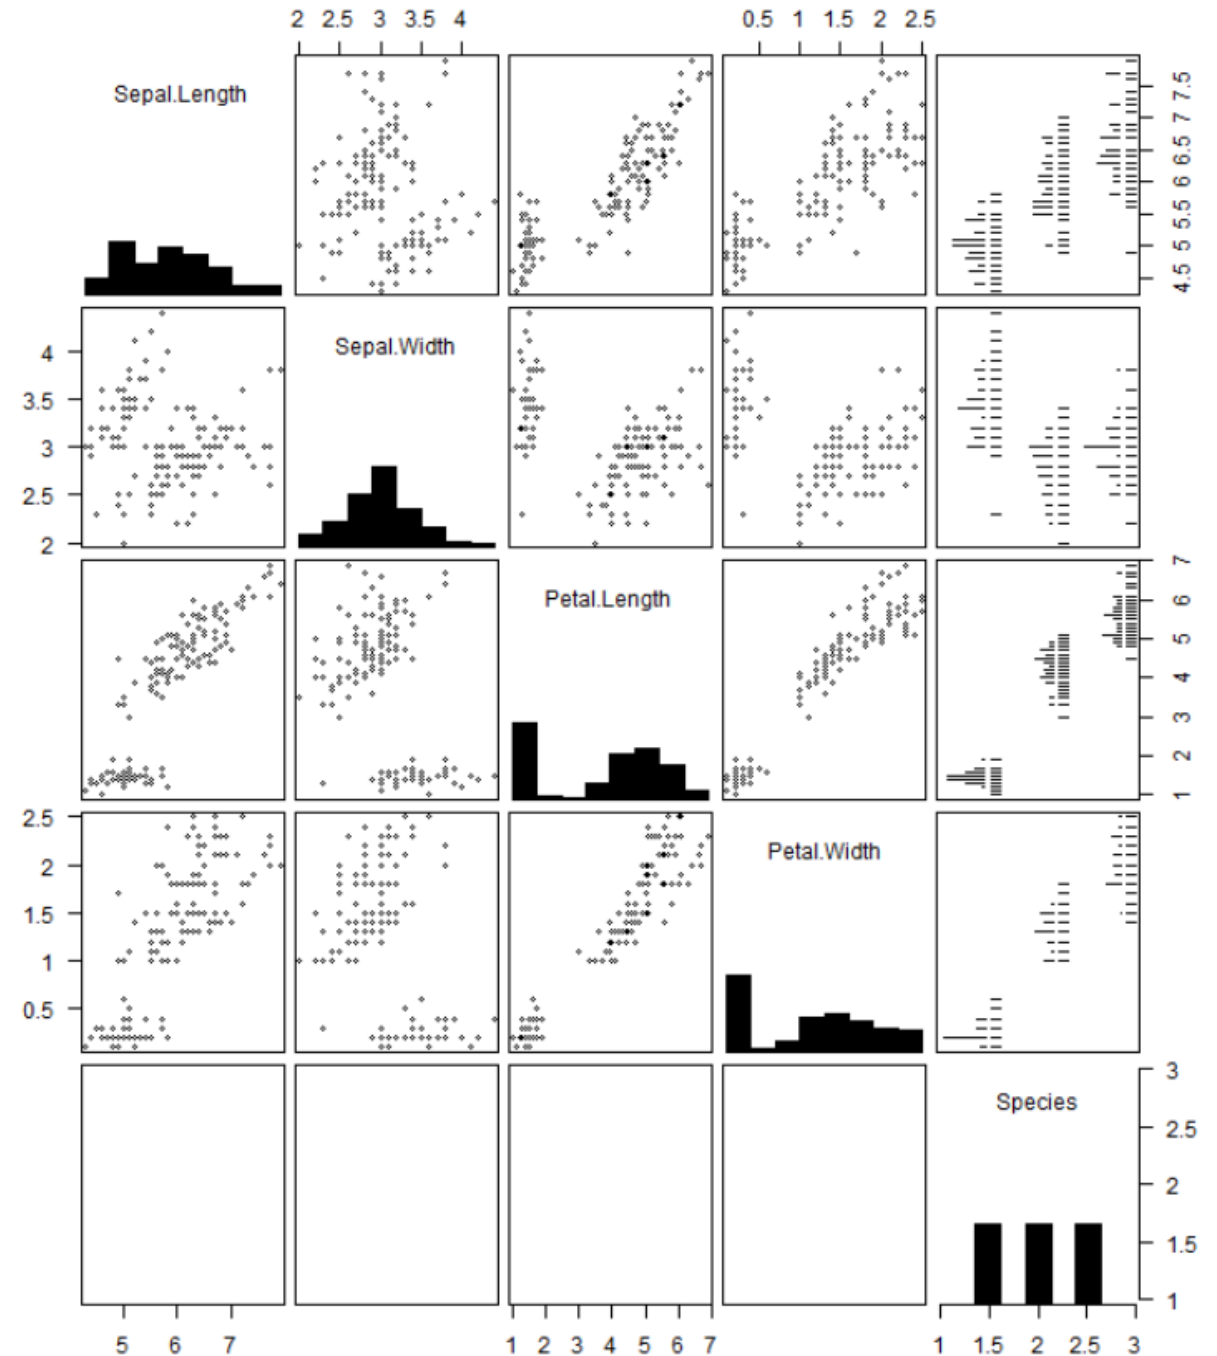
```
library (car)
```

```
scatterplotMatrix (~  
  Sepal.Length + Sepal.Width +  
  Petal.Length + Petal.Width |  
  Species, data=iris,  
  diagonal="histogram")
```




Visualizing Relationships: Continuous & Discrete

```
install.packages("gpairs")  
library(gpairs)  
gpairs(iris[, 1:5])  
help(gpairs)
```




Visualizing Correlation Matrix

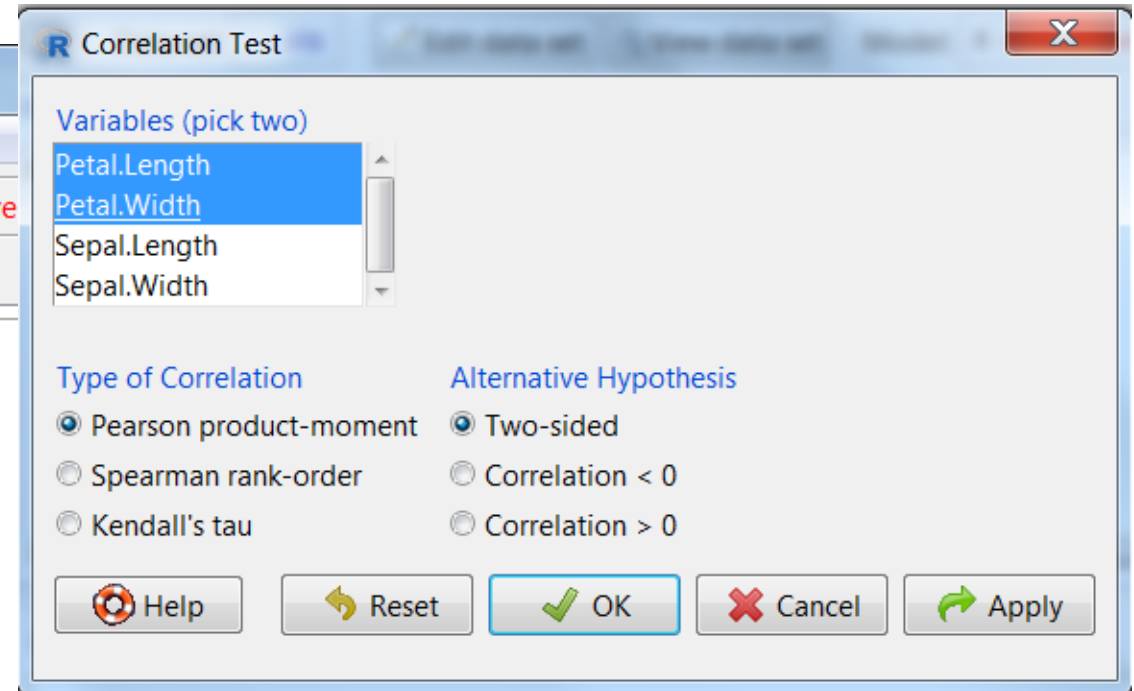
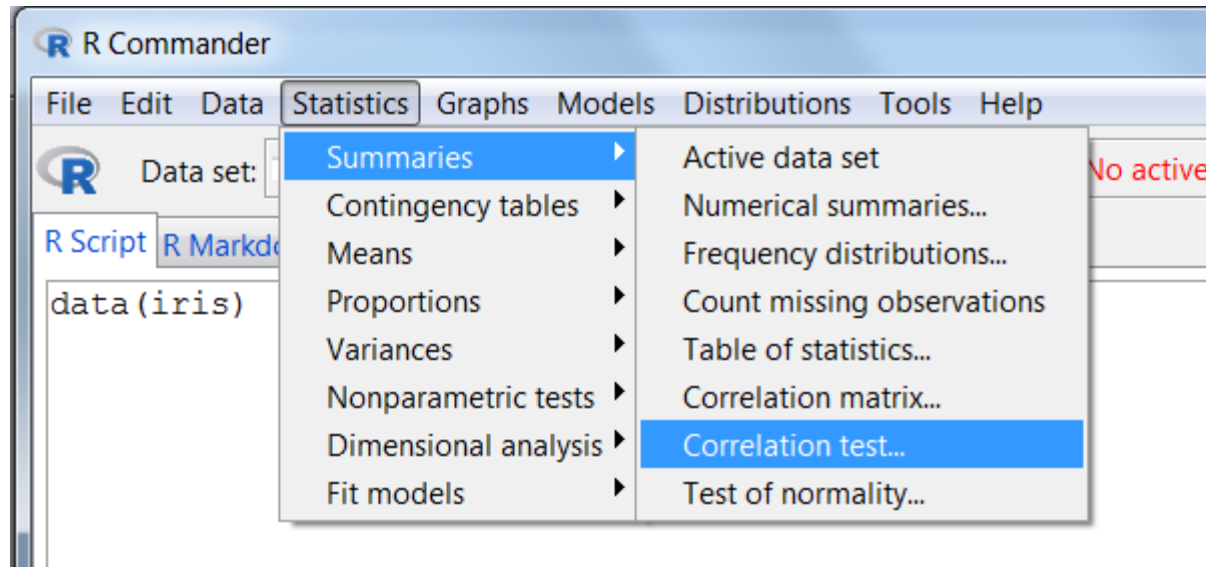
```
install.packages ("corrplot")  
library (corrplot)  
corrplot (cor (iris[ , 1:4], order="hclust")  
help (corrplot)
```



```
library (corrplot)  
library(gplots)  
corrplot.mixed (cor (iris[ , 1:4]),  
upper="ellipse", tl.pos="lt",  
col=colorpanel(50, "red", "grey60", "blue4"))
```



Statistical Significance of Correlations



```
Pearson's product-moment correlation  
  
data: Petal.Length and Petal.Width  
t = 43.387, df = 148, p-value < 2.2e-16  
alternative hypothesis: true correlation is not equal to 0  
95 percent confidence interval:  
 0.9490525 0.9729853  
sample estimates:  
      cor  
0.9628654
```

$p\text{-value} < 0.05$: statistically significant

Linear Transformation Matrix

SCALING, REFLECTION, ROTATION

Linear Transformation Matrix

Matrix-based Transformation of Vectors

$$V_{old} \xrightarrow{\text{matrix, } A} V_{new}$$

via **vector-matrix multiplication**

- **Scaling Matrix**
- **Reflection Matrix**
- **Rotation Matrix**
- **Projection Matrix**

Notation

$$\mathbf{u} = [u_1, u_2, \dots, u_p] - \text{row vector}$$

$$\mathbf{v} = \begin{bmatrix} v_1 \\ \dots \\ v_m \end{bmatrix} - \text{column vector}$$

$$\mathbf{A} = \mathbf{A}_{m \times p} = \begin{bmatrix} a_{11} & \dots & a_{1p} \\ \dots & \dots & \dots \\ a_{m1} & \dots & a_{mp} \end{bmatrix} - m \times p \text{ matrix}$$

$$\mathbf{I} = \mathbf{I}_{m \times m} = \begin{bmatrix} 1 & 0 & \dots & 0 \\ 0 & 1 & \dots & 0 \\ \dots & \dots & \dots & \dots \\ 0 & 0 & 0 & 1 \end{bmatrix} - \text{identity matrix}$$

Scaling Matrix


Row vectors:

$$v_{old} \in R^p \xrightarrow[\text{matrix, } A]{\text{scalar, } \lambda \in R^+} v_{new} = \lambda \cdot v_{old} \in R^p$$

- **Unchanged:**
 - **Direction** of a vector if $\lambda > 0$
- **Changed:**
 - **Vector norm/length** $\|v_{new}\| = \lambda \cdot \|v_{old}\|$
 - **Direction** of a vector if $\lambda < 0$

$$A_\lambda = \lambda \cdot I_{p \times p} = \begin{bmatrix} \lambda & 0 & \dots & 0 \\ 0 & \lambda & \dots & 0 \\ \dots & \dots & \dots & \dots \\ 0 & 0 & 0 & \lambda \end{bmatrix} - \text{Scaling Matrix}$$

$$v_{new} = v_{old} \times A_\lambda = v_{old} \cdot \lambda \cdot I_{p \times p}$$



Reflection Matrix

Row vectors:

$$v_{old} \in R^p \xrightarrow{\text{matrix, } A} v_{new} = v_{old} \cdot A \in R^p$$


$$A = I'_{p \times p} = \begin{bmatrix} 1 & 0 & \dots & 0 \\ 0 & -1 & \dots & 0 \\ \dots & \dots & \dots & \dots \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Reflection Matrix: Identity Matrix with some diagonal elements set to -1

$$v_{new} = v_{old} \cdot A$$

Equivalent to: multiplying one or more vector components by -1

- **Changed:**
 - **Direction** of a vector
 - Reflects across one/more axis
- **Unchanged:**
 - Vector **norm/length**



Rotation Matrix

Row vectors:

$$v_{old} \in R^p \xrightarrow[\substack{\text{matrix, } A \\ \det(A)=1 \\ A^T=A^{-1}}]{ } v_{new} = v_{old} \cdot A \in R^p$$

- **Changed:**
 - **Direction** of a vector
- **Unchanged:**
 - **Vector norm/length**


$$A = A_{p \times p} = \begin{bmatrix} a_{11} & \dots & a_{1p} \\ \dots & \dots & \dots \\ a_{p1} & \dots & a_{pp} \end{bmatrix} - \text{Rotation Matrix: } p \times p \text{ orthogonal matrix}$$

$$A^T = A^{-1} - \text{orthogonal matrix}$$

$$\det(A)=1$$

Example: $p = 2$

$$A(\theta) = \begin{bmatrix} \cos \theta & \sin \theta \\ -\sin \theta & \cos \theta \end{bmatrix} - \theta\text{-angle rotation: clockwise } (A \times v_{old}) \text{ or counterclockwise } (v_{old} \times A)$$



Linear Transformation Matrix

PROJECTION

Projection Matrix


Row vectors:

$$v_{old} \in \mathbb{R}^p \longrightarrow v_{new} = v_{old} \times A_{p \times d} \in \mathbb{R}^d, d < p$$

- **Changed:**
 - **Dimensionality** of a vector
 - **Direction** of a vector
 - Vector **norm/length**


Example 1: $p=2 \rightarrow d=1$

Original data, $p=2$



$$X_{4 \times 2}$$


x_1 -projection,



“loss of info”

$$X'_{4 \times 2} = X_{4 \times 2} \cdot \begin{bmatrix} 1 & 0 \\ 0 & 0 \end{bmatrix}$$

x_2 -projection, $k=1$



“no loss”


$$X'_{4 \times 2} = X_{4 \times 2} \cdot \begin{bmatrix} 0 & 0 \\ 0 & 1 \end{bmatrix}$$

$$X'_{m \times d} = X_{m \times d} \cdot P_{d \times d}$$


Which projection is “better”?

$$P_{d \times d} = \begin{bmatrix} 1 & 0 & \dots & 0 \\ 0 & 0 & \dots & 0 \\ \dots & \dots & \dots & \dots \\ 0 & 0 & 0 & 1 \end{bmatrix} - \text{projection matrix; some diagonal elements are 0}$$

Example 2: Linear, Orthogonal Projection



Is there a “better” projection?





Projection:

- **Linear**, e_1 – line
- **Orthogonal**

$$e_1 \perp e_2$$

e_1, e_2 – **eigenvectors** of X


Example 3: Non-linear projection



NON-linear projection

Example 4: Projection for Labeled Data

What is a “better” projection?



Ideal for machine learning (ML) algorithm:

- **Points within the same group come from a Gaussian distribution (spherical shapes)**
- **Points from different groups are linearly separable**

Summary: Types of Dimension Reduction

- **Linear vs. NON-linear**
- **Orthogonal vs. non-orthogonal**
- **Unsupervised (unlabeled data) vs. supervised (labeled data)**

Linear dimension reduction:

- Interpretable in original space
- Preserves non-linearity for visualization
- Orthogonal projections
- Non-orthogonal projections
- Favored for structure discovery

Non-linear dimension reduction:

- Lower-dimensional representation
- Interpretable w.r.t. Non-linear transformation
- 1 to 3 orders of magnitude more computation
- Favored for prediction or classification

- **A lower-dimensional representation that contains the essence of the high dimensional data**
- **Blessing of dependence/correlation that saves us from the curse of dimensionality**

Linear Orthogonal Projection

PCA: PRINCIPAL COMPONENT ANALYSIS

What is “better” projection: $d \rightarrow k$ ($k < d$)?

- Many definitions are possible
- Definition 1:
 - Projection that **maximizes the VARIANCE** of the original d -dimensional data upon its projection onto the target k -dimensions ($k < d$)

R Example: PCA with prcomp()

```
data (iris)
```

```
View (iris)
```

```
X <- iris [, 1:4]
```

```
plot (X[,1], X[,2], pch=19)
```

```
points (X[1:50, 1], X[1:50,2], col="red", pch=19)
```

```
points (X[51:100, 1], X[51:100,2], col="green", pch=19)
```

```
pca.model <- prcomp (X)
```

```
pca.model$sdev           # square roots of eigenvalues covariance/correlation matrix
```

```
plot (pca.model, type="l") # screeplot
```

```
pca.model$rotation       # matrix of eigenvector columns
```

```
pca.model
```

```
summary (pca.model) # check proportion of variance
```

```
P <- pca.model$x         # projection of X onto eigenvectors
```



```
dim (P)
```

```
head (P)
```

```
plot (P[,1], P[,2], pch=19)
```

```
points (P[1:50, 1], P[1:50,2], col="red", pch=19)
```


```
points (P[51:100, 1], P[51:100,2], col="green", pch=19)
```



Visualizing PCA results

Screenplot

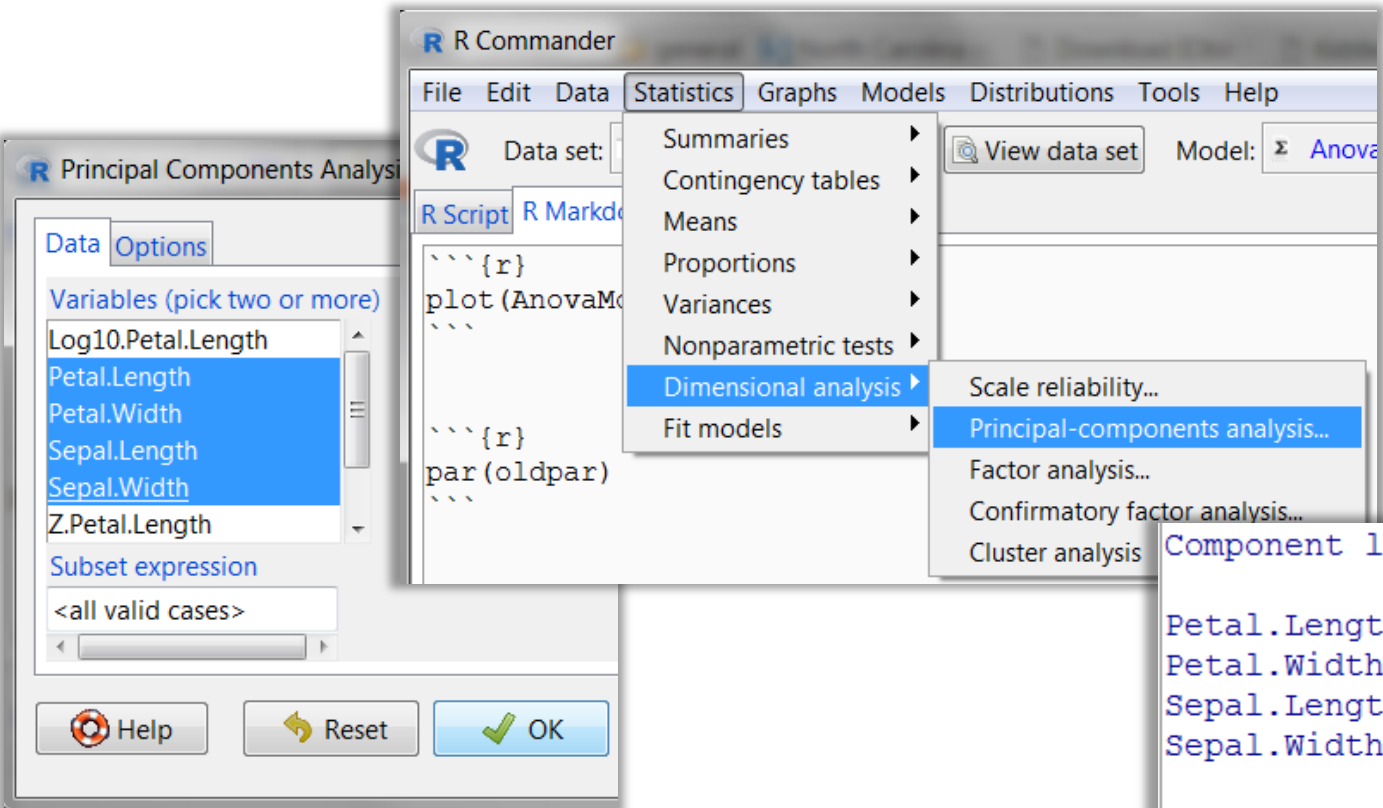
`plot(pca.model, type="l")`



Successive variance accounted
by each component

`biplot(pca.model)`

Principal Component Analysis with R Commander



The screenshot shows the R Commander interface. The 'Statistics' menu is open, and the path 'Dimensional analysis' > 'Principal-components analysis...' is highlighted. To the left, the 'Principal Components Analysis' dialog box is visible, showing selected variables: Log10.Petal.Length, Petal.Length, Petal.Width, Sepal.Length, Sepal.Width, and Z.Petal.Length. The 'Subset expression' is set to '<all valid cases>'. Below the dialog, the output of the PCA is displayed in a text box.

Component loadings:

	Comp.1	Comp.2	Comp.3	Comp.4
Petal.Length	0.5804131	0.02449161	0.1421264	0.8014492
Petal.Width	0.5648565	0.06694199	0.6342727	-0.5235971
Sepal.Length	0.5210659	0.37741762	-0.7195664	-0.2612863
Sepal.Width	-0.2693474	0.92329566	0.2443818	0.1235096

Component variances:

	Comp.1	Comp.2	Comp.3	Comp.4
2.91849782	0.91403047	0.14675688	0.02071484	

Importance of components:

	Comp.1	Comp.2	Comp.3	Comp.4
Standard deviation	1.7083611	0.9560494	0.38308860	0.143926497
Proportion of Variance	0.7296245	0.2285076	0.03668922	0.005178709
Cumulative Proportion	0.7296245	0.9581321	0.99482129	1.000000000

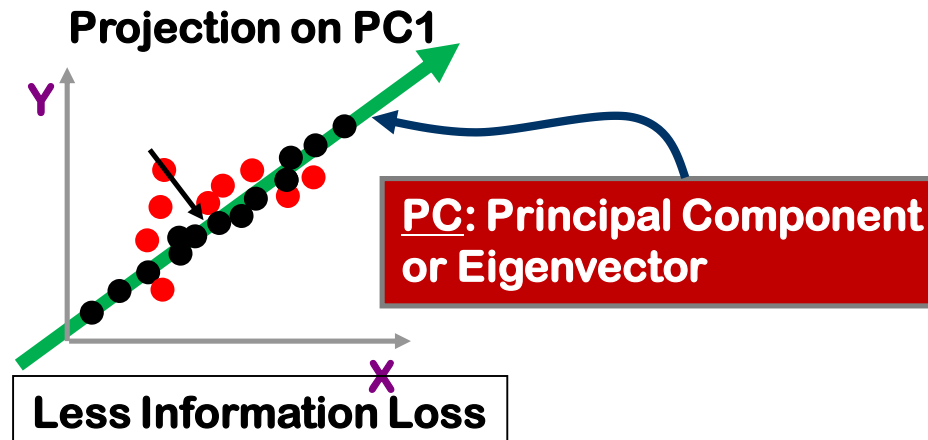
An arrow points to the value 0.9581321 in the 'Cumulative Proportion' row, which is circled in red.

PCA: Linear Orthogonal Dimension Reduction

Principal Component Analysis (**PCA**) finds **intrinsic** dimensionality and allows for low-dimensional representation.



Principal Component Analysis is the **Spectral Decomposition of the Covariance Matrix**.



Matrix Decomposition

PCA: SVD OF COVARIANCE MATRIX

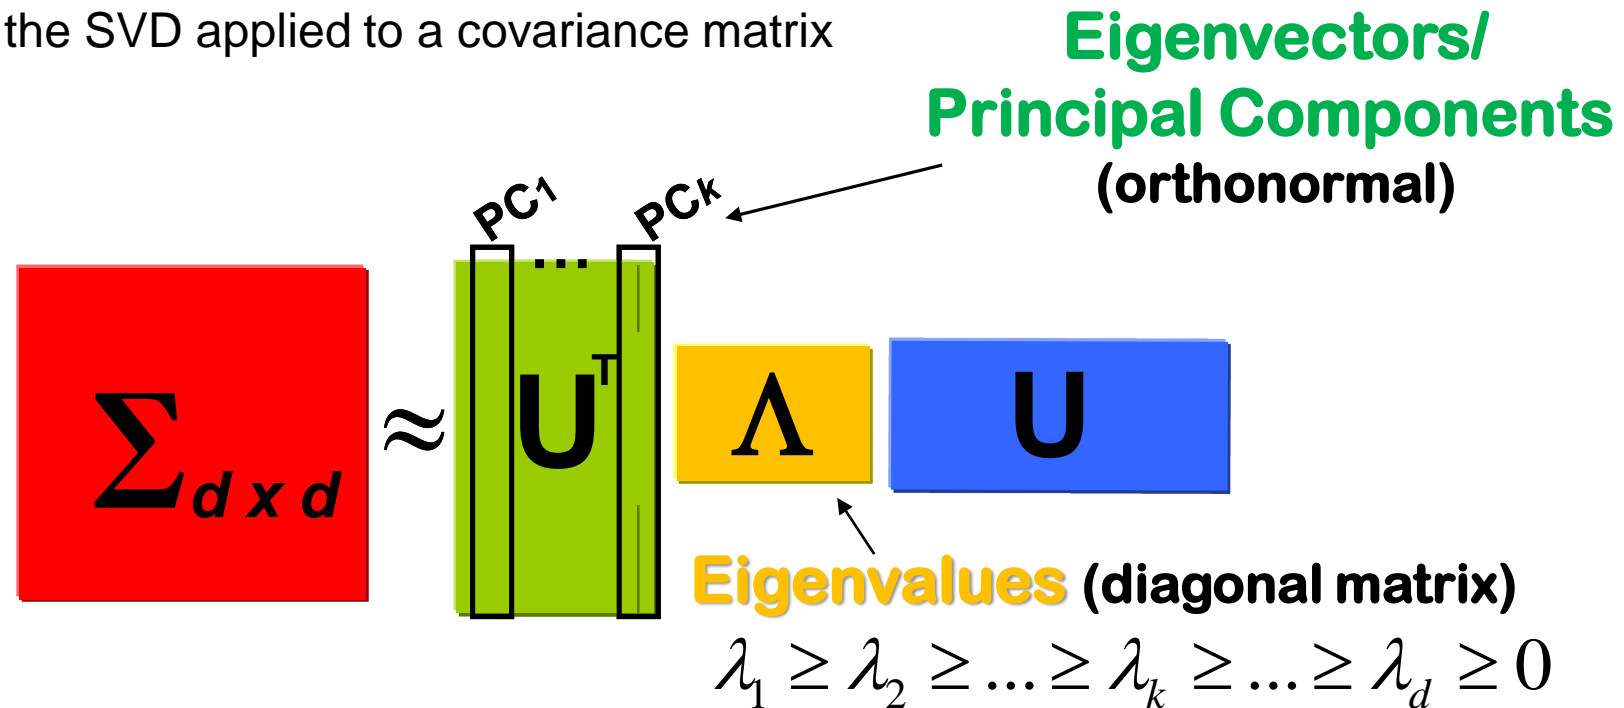
PCA is the SVD of Covariance/Correlation Matrix

Singular Value Decomposition (SVD)

The technique underlying PCA analysis

PCA = SVD (Covariance Matrix)

PCA is the SVD applied to a covariance matrix




Details

`princomp` is a generic function with "formula" and "default" methods.

The calculation is done using [eigen](#) on the correlation or covariance matrix, as determined by [cor](#).

Extracted Principal Components (PCs) are Uncorrelated

```
> cor (pca.model$x)
      PC1      PC2      PC3      PC4
PC1  1e+00 -1e-16 -1e-15  2e-15
PC2 -1e-16  1e+00  2e-16 -5e-16
PC3 -1e-15  2e-16  1e+00 -1e-15
PC4  2e-15 -5e-16 -1e-15  1e+00
> corrplot (cor (pca.model$x))
```



Preserved Variability for Top-k PCs

Percentage of variability preserved
if the first k PCs are used for projection:

$$\frac{\sum_{i=1}^k \lambda_i}{\sum_{i=1}^d \lambda_i} = \frac{\sum_{i=1}^k \lambda_i}{\text{trace}(\Sigma)}$$

SORTED Eigenvalues


$$\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_k \geq \dots \geq \lambda_d \geq 0$$




$$\text{Preserved variability} = \frac{\text{Sum of } k \text{ eigenvalues}}{\text{Sum of all eigenvalues}}$$

with k dimensions

covariance
matrix



≈



Eigenvalues (diagonal matrix)

Key Points about PCA

- PCA = SVD (Covariance Matrix, Σ)
- Linear, orthogonal projection:
 - “Best” linear orthogonal k -dimensional ($k < d$) view of data
- How “good” the k -dimensional view is:

$$\frac{\sum_{i=1}^k \lambda_i}{\sum_{i=1}^d \lambda_i} = \frac{\sum_{i=1}^k \lambda_i}{\text{trace}(\Sigma)}$$

Eigenvalues: Importance of eigenvectors

$$\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_k \geq \dots \geq \lambda_d \geq 0$$

Importance of Principal Components (PC), or Eigenvectors, or Rotations

PC1 preserves more variance than PC2


PC2 preserves more variance than PC3

....

Proportion of Variance Preserved if only k PCs are used:

$$\frac{\sum_{i=1}^k \lambda_i}{\sum_{i=1}^d \lambda_i} = \frac{\sum_{i=1}^k \lambda_i}{\text{trace}(\Sigma)}$$

Eigenvectors are NOT unique



PCA: Covariance vs. Correlation Matrix

- PCA results **depend on the scales** at which the variables are measured:
 - PCA should only be used with the raw data if all variables have the same units of measure
- PCA results **depend on the variances** of the variables: the ones with the highest sample variances will tend to be emphasized in the first few principal components:
 - Use PCA with covariance matrix only if you wish to give variables with higher variances more weight in the analysis
- If the variables either have **different units of measurement** (i.e., pounds, feet, gallons, etc), or if we wish **each variable to receive equal weight in the analysis**, then the variables should be **standardized (Z-scores)** before a principal components analysis is carried out.

Matrix Decomposition

PCA: DIMENSION REDUCTION VIA FEATURE EXTRACTION

PC is a weighted linear sum of original features

Extracted Feature → PCA-based Feature Extraction:

$$PC = w_1 * f_1 + w_2 * f_2 + \dots + w_d * f_d$$

The magnitude of each weight indicates how important the corresponding feature is → it could be used as a *feature selection* technique!

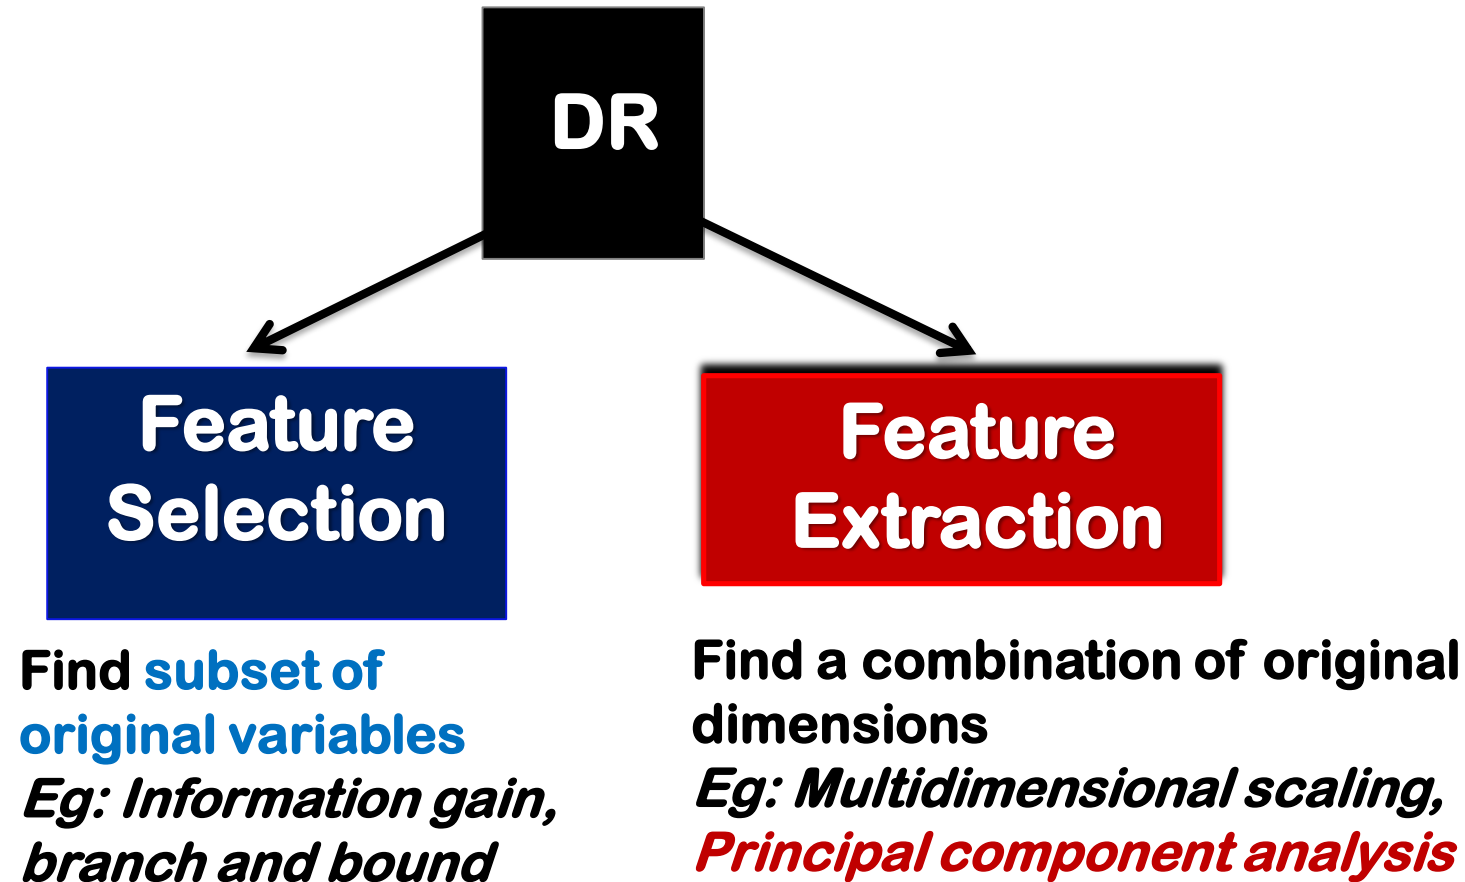
```
> options(digits=1)
> pca.model$rotation
```

	PC1	PC2	PC3	PC4
Sepal.Length	0.36	-0.66	0.58	0.3
Sepal.Width	-0.08	-0.73	-0.60	-0.3
Petal.Length	0.86	0.17	-0.08	-0.5
Petal.Width	0.36	0.08	-0.55	0.8

PC1 = 0.86 Petal.Length + 0.36 Petal.Width
+ 0.36 Sepal.Length – 0.08 Sepal.Width

PC2 = -0.73 Sepal.Width - 0.66 Sepal.Length
+ 0.17 Petal.Length + 0.08 Petal.Width

Classification of Dimension Reduction (DR) Methods



Motivation for Dimension Reduction

- Not all the measured variables are important for understanding the underlying “interesting” phenomena – could complicate the process of data analysis
- Decrease the **computational cost** for other data mining tasks:
 - Proximity measure calculations: $O(d) \rightarrow O(k)$
- Reduce the **noise** in the data
- Improve the **accuracy** of predictive models
- Reduce **collinearity** among variables/features

What is Dimensionality Reduction?

Objective: To transform data from a **high-dimensional** space to a corresponding representation in some **low-dimensional** space, while “best” preserving the information.

Given dataset with m objects:

$$\begin{array}{ccc} X \in \mathbb{R}^{m \times n} & \xrightarrow{\text{DR}} & Y \in \mathbb{R}^{m \times p} \\ m \text{ dimensions} & & p \text{ dimensions} \end{array}$$

where $p \ll n$