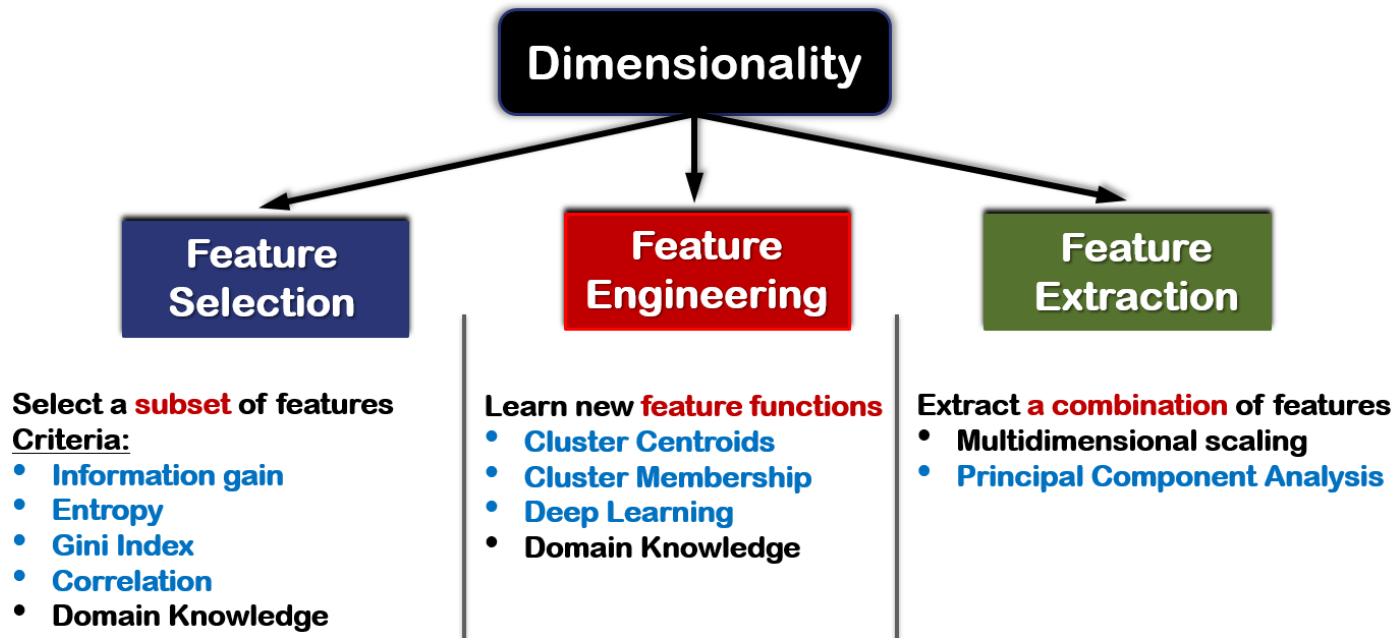


PCA: Dimension Reduction



- Dimensionality and underdetermined problems
- Feature selection vs. feature extraction vs. feature engineering
- Dimension reduction (DR): unsupervised vs. supervised; linear vs. non-linear; orthogonal vs. non-orthogonal
- DR: linear, orthogonal with Principal Component Analysis (PCA)
- Eigenvalue and eigenvectors


Prof. Nagiza F. Samatova

samatova@csc.ncsu.edu

Department of Computer Science
North Carolina State University

Dimension Reduction

MOTIVATION





What Analysis Methods to Use?

Analysis methods fail for a few **gigabytes**.

Method Complexity:

Calculate means $O(n)$

Calculate Histogram $O(n \log(n))$

Calculate PCA $O(n \cdot d)$

Clustering algorithms $O(n^2)$

If $n = 10\text{GB}$, then what is $O(n)$ or $O(n^2)$ on a teraflop computers?

$1\text{GB} = 10^9$ bytes $1\text{Tflop} = 10^{12}$ op/sec

$O(\dots)$ - in the order of

n – number of rows

d - number of columns

Data size n	Algorithm Complexity		
	n	$n \log(n)$	n^2
100B	10^{-10} sec.	10^{-10} sec.	10^{-8} sec.
10KB	10^{-8} sec.	10^{-8} sec.	10^{-4} sec.
1MB	10^{-6} sec.	10^{-5} sec.	1 sec.
100MB	10^{-4} sec.	10^{-3} sec.	3 hrs
10GB	10^{-2} sec.	0.1 sec.	3 yrs.


For illustration chart assumes 10^{-12} sec.
(**1Tflop/sec**) calculation time per data point

How to Make Sense of Data?

Know Your Limits & Be Smart!



Not humanly possible to browse a petabyte of data.
Analysis must reduce data to quantities of interest.



Computations:

Must be smart about which probe combinations to see!


Physical Experiments:

Must be smart about probe placement!

To see 1 percent of a petabyte at 10 megabytes per second takes:

35 8-hour days!

It is not just the Size



High-dimensional Data

- **Text Data**

- Record/Row: Document ID
- Dimension/Column: Each word in a collection of text documents

- **Image Data**

- Record/Row: Image ID
- Dimension/Column: Each pixel

- **Audio Data**

- Record/Row: Audio record ID
- Dimension/Column: Frequency of an audio signal OR a word if audio → text conversion

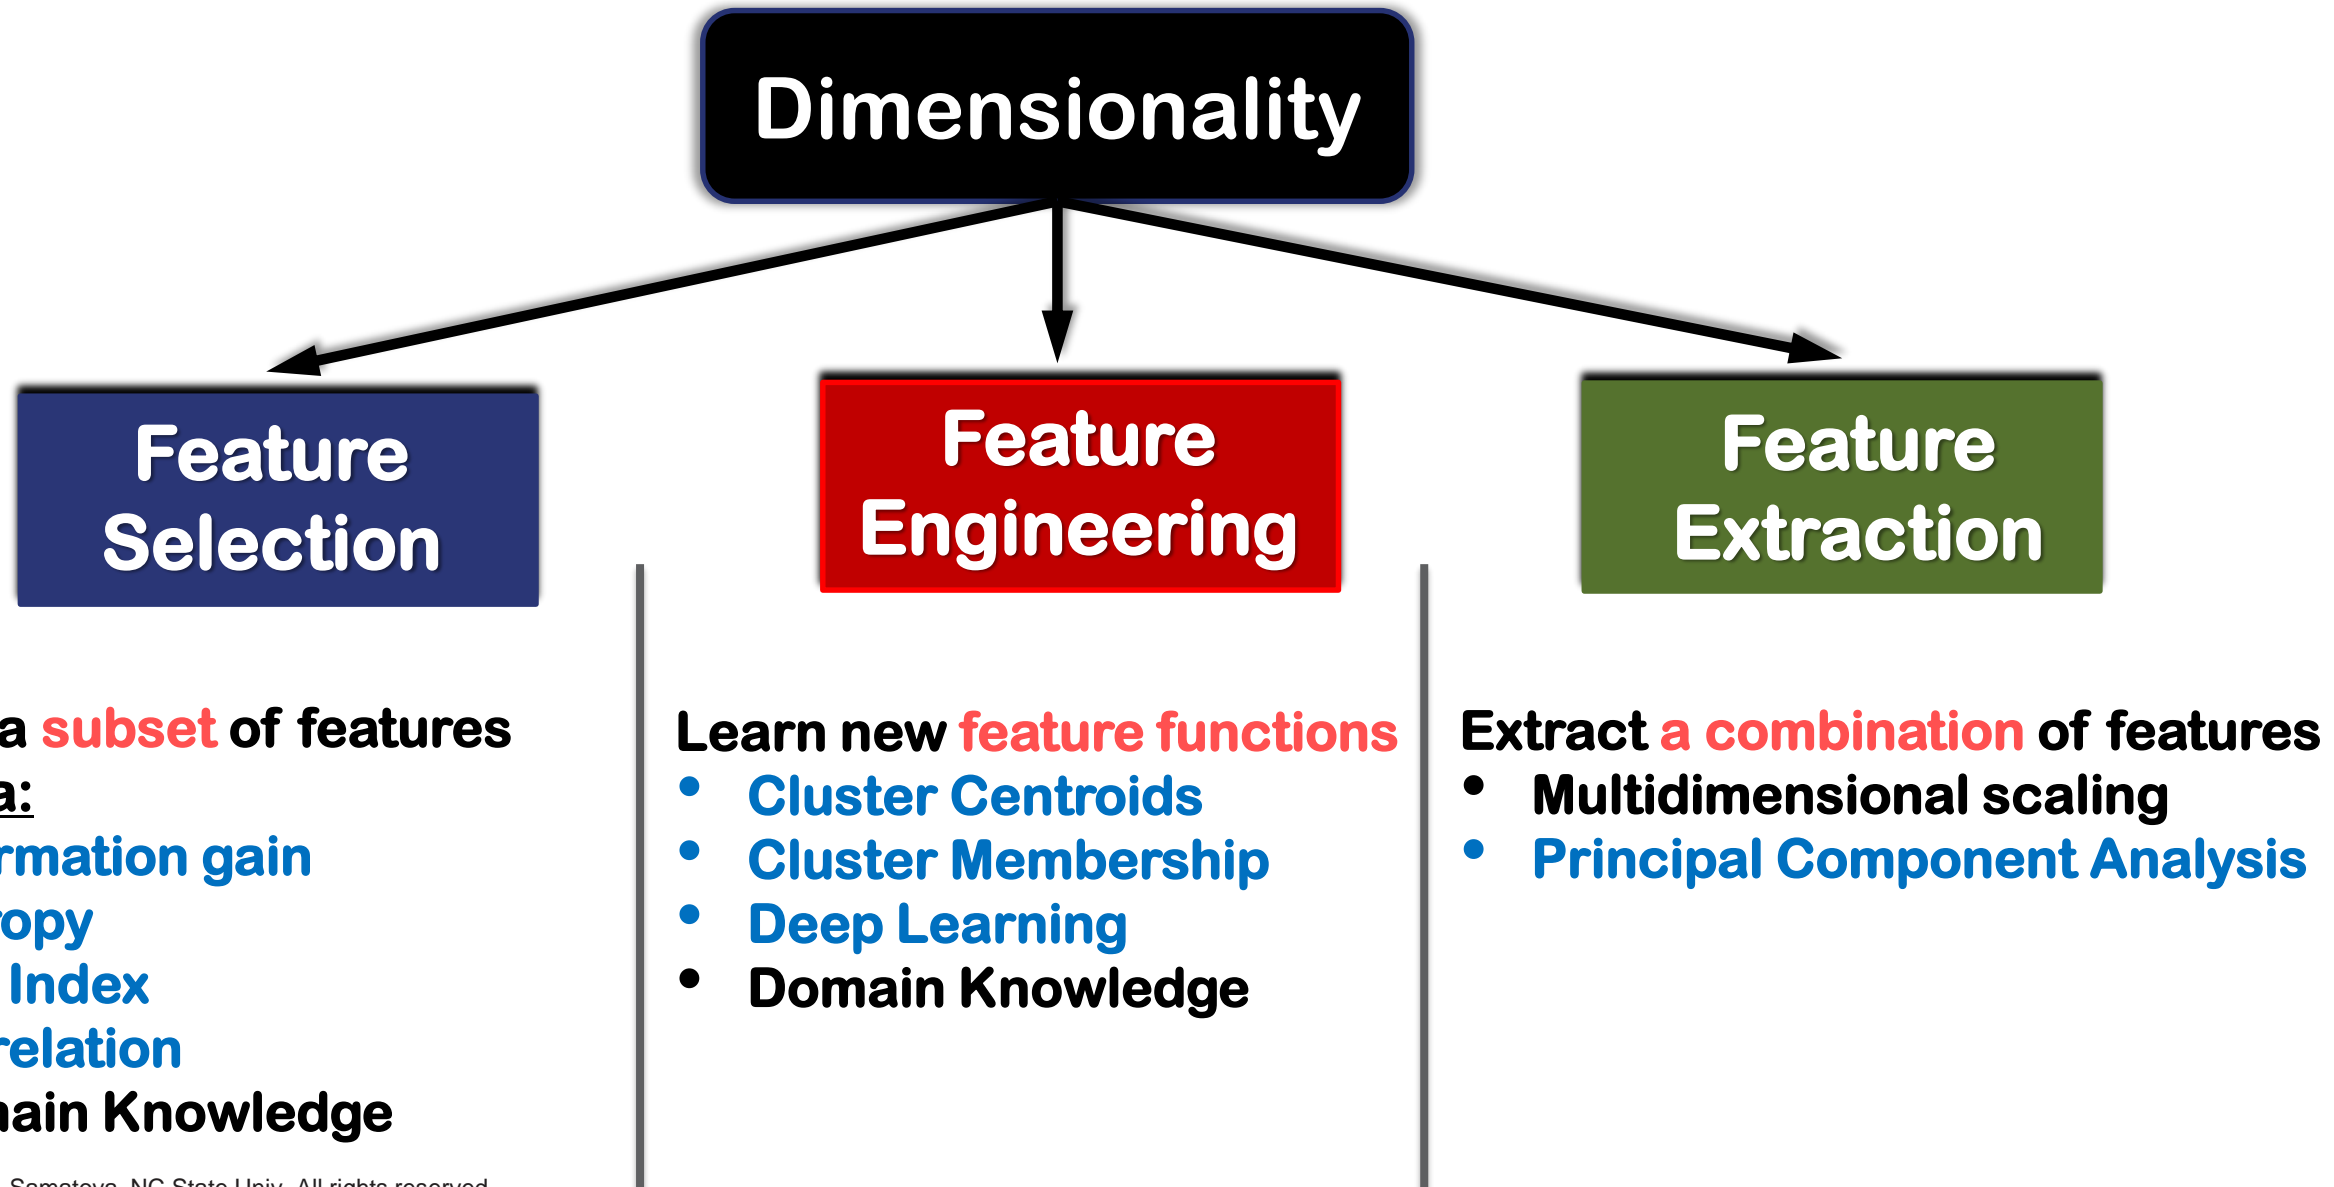
- **Sales Data**

- Record/Row: Transaction ID
- Dimension/Column: Each Product ID

Taxonomy: Dimensionality-aware Methods

**FEATURE SELECTION, FEATURE
EXTRACTION, FEATURE ENGINEERING**

Dimensionality Challenge: Strategies to Cope With



Motivation: Why to Cope with Dimensionality Issue?

- Not all the measured variables are important for understanding the underlying “interesting” phenomena:
 - could complicate the process of data analysis
- Decrease the **computational cost** for other data mining tasks:
 - Proximity measure calculations: $O(d) \rightarrow O(k)$, $d \gg k$
- Reduce the **noise** in the data:
 - improve signal to noise ratio
- Improve the **accuracy** of predictive models
 - better designed/engineered features capable of capturing non-linear signals in the data
- Reduce **collinearity** among variables/features
 - Critical for regression models

What is Dimensionality Reduction (DR)?

Objective: To transform data from a **high-dimensional** representation to a **low-dimensional** representation, while “**best**” preserving the information.

Given dataset with n objects:

$$\begin{array}{ccc} X \in \mathbb{R}^{n \times m} & \xrightarrow{\text{DR}} & X' \in \mathbb{R}^{n \times p} \\ m \text{ dimensions} & & p \text{ dimensions} \end{array}$$


where $p \ll m$

Text Mining Example:

- $m = 50,000$ words $\rightarrow p \approx 100$ extracted features


Example 1: $p=2 \rightarrow d=1$

Original data, $p = 2$




$$X_{4 \times 2}$$

x_1 -projection, $k=1$



$$X'_{4 \times 2} = X_{4 \times 2} \cdot \begin{bmatrix} 1 & 0 \\ 0 & 0 \end{bmatrix}$$

x_2 -projection, $k = 1$



$$X'_{4 \times 2} = X_{4 \times 2} \cdot \begin{bmatrix} 0 & 0 \\ 0 & 1 \end{bmatrix}$$


$$X'_{m \times d} = X_{m \times d} \cdot P_{d \times d}$$

$$P_{d \times d} = \begin{bmatrix} 1 & 0 & \dots & 0 \\ 0 & 0 & \dots & 0 \\ \dots & \dots & \dots & \dots \\ 0 & 0 & 0 & 1 \end{bmatrix} - \text{projection matrix; some diagonal elements are 0}$$


Which projection is **“better”**?

Example 2: Linear, Orthogonal Projection


Original data, $d=2$



x_1 -projection, $k=1$




x_2 -projection, $k=1$




Is there a “better” projection?

Another basis, $d=2$
(rotate coord. system)



e_1 -projection,
 $k=1$





Projection:

- Linear, e_1 – line
- Orthogonal

$$e_1 \perp e_2$$

e_1, e_2 – eigenvectors of X


Example 3: Non-linear projection



NON-linear projection

Example 4: Projection for Labeled Data

What is a “better” projection?



Ideal for machine learning (ML) algorithm:

- Points within the same group come from a Gaussian distribution (spherical shapes)
- Points from different groups are linearly separable

“better” means:

- it bests, ideally, **linearly separates** different groups of data, i.e.
- points from **the same group** are **closer** to each other and are **farther away** from the points in **different groups**

Summary: Taxonomy of Dimension Reduction

- **Linear vs. NON-linear**
- **Orthogonal vs. non-orthogonal**
- **Unsupervised (unlabeled data) vs. supervised (labeled data)**

Linear dimension reduction:

- Interpretable in original space
- Preserves non-linearity for visualization
- Orthogonal projections
- Non-orthogonal projections
- Favored for structure discovery

Non-linear dimension reduction:

- Lower-dimensional representation
- Interpretable w.r.t. Non-linear transformation
- 1 to 3 orders of magnitude more computation
- Favored for prediction or classification

- **A lower-dimensional representation that contains the essence of the high dimensional data**
- **Blessing of dependence/correlation that saves us from the curse of dimensionality**

Linear Orthogonal Dimension Reduction

PCA: PRINCIPAL COMPONENT ANALYSIS

Dimensionality Challenge: Strategies to Cope With

Dimensionality

```
graph TD; Dimensionality --> FeatureSelection[Feature Selection]; Dimensionality --> FeatureEngineering[Feature Engineering]; Dimensionality --> FeatureExtraction[Feature Extraction];
```

Feature Selection

Select a **subset** of features

Criteria:

- Information gain
- Entropy
- Gini Index
- Correlation
- Domain Knowledge

Feature Engineering

Learn new **feature functions**

- Cluster Centroids
- Cluster Membership
- Deep Learning
- Domain Knowledge

Feature Extraction

Extract a **combination** of features

- Multidimensional scaling
- Principal Component Analysis

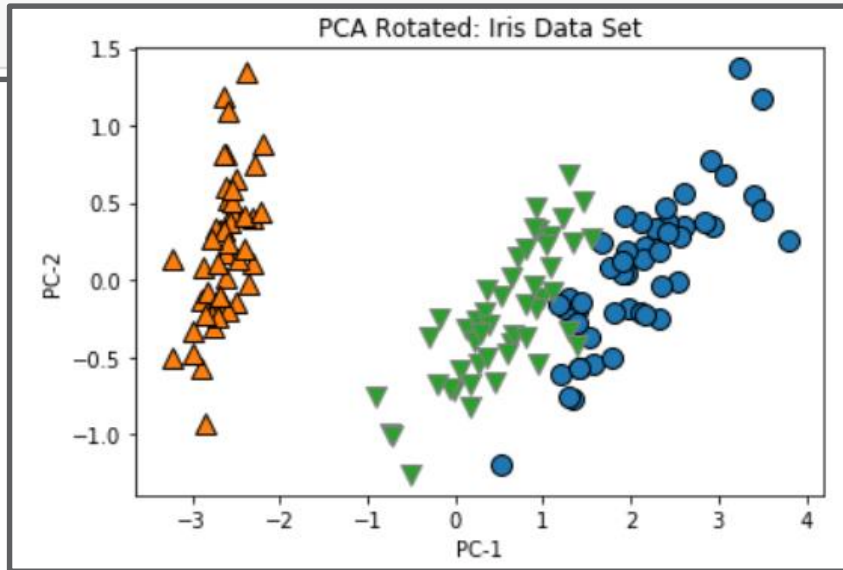
What is “better” projection: $d \rightarrow k$ ($k < d$)?

- Many definitions are possible
- Definition 1:
 - Projection that **maximizes the VARIANCE** of the original d -dimensional data upon its projection onto the target k -dimensions ($k < d$)

Python Code Example: Iris Data

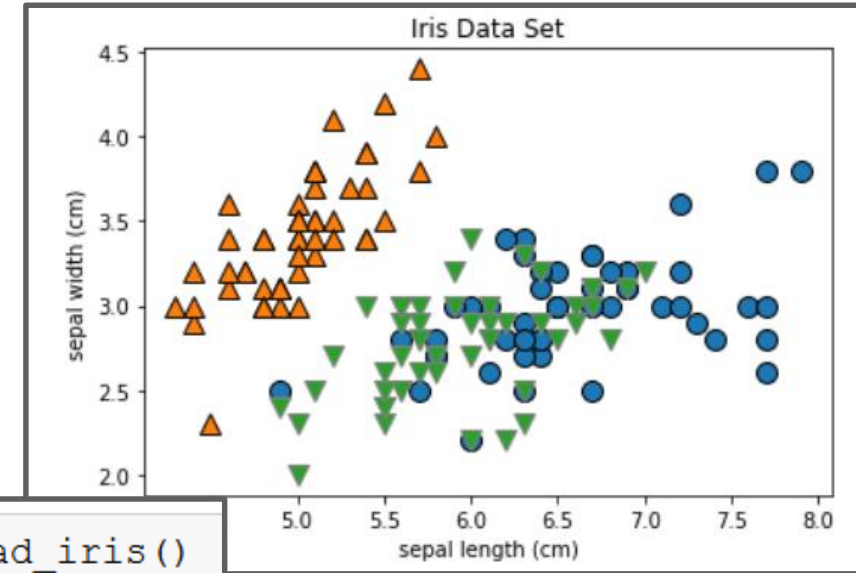
```
pca = decomposition.PCA(n_components=3)
pca.fit(X)
X_rot = pca.transform(X)
```

```
# Plot rotated 4-D iris data in 2D: the first two PCs
mglearn.discrete_scatter(X_rot[:, 0], X_rot[:, 1], y)
plt.title("PCA Rotated: Iris Data Set")
plt.xlabel("PC-1")
plt.ylabel("PC-2")
plt.show()
```



PCA-transformed data in 2-d

original 4-d data in 2-d

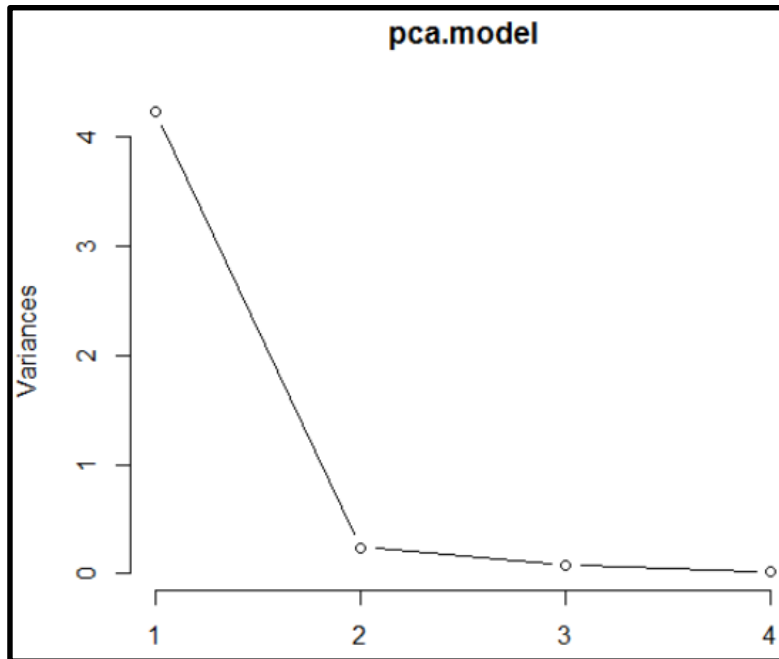


```
iris = datasets.load_iris()
X = iris.data
y = iris.target
print(iris.feature_names)
print(X[0:5])
print(iris.target_names)
print(y[[0, 51, 101]])
```

```
# Plot original 4-D iris data in 2D
mglearn.discrete_scatter(X[:, 0], X[:, 1], y)
plt.title("Iris Data Set")
plt.xlabel("sepal length (cm)")
plt.ylabel("sepal width (cm)")
plt.show()
```

Visualizing PCA Results: Iris Data

Screeplot



Successive variance accounted
by each component

Component loadings:

	Comp.1	Comp.2	Comp.3	Comp.4
Petal.Length	0.5804131	0.02449161	0.1421264	0.8014492
Petal.Width	0.5648565	0.06694199	0.6342727	-0.5235971
Sepal.Length	0.5210659	0.37741762	-0.7195664	-0.2612863
Sepal.Width	-0.2693474	0.92329566	0.2443818	0.1235096

Component variances:

	Comp.1	Comp.2	Comp.3	Comp.4
	2.91849782	0.91403047	0.14675688	0.02071484

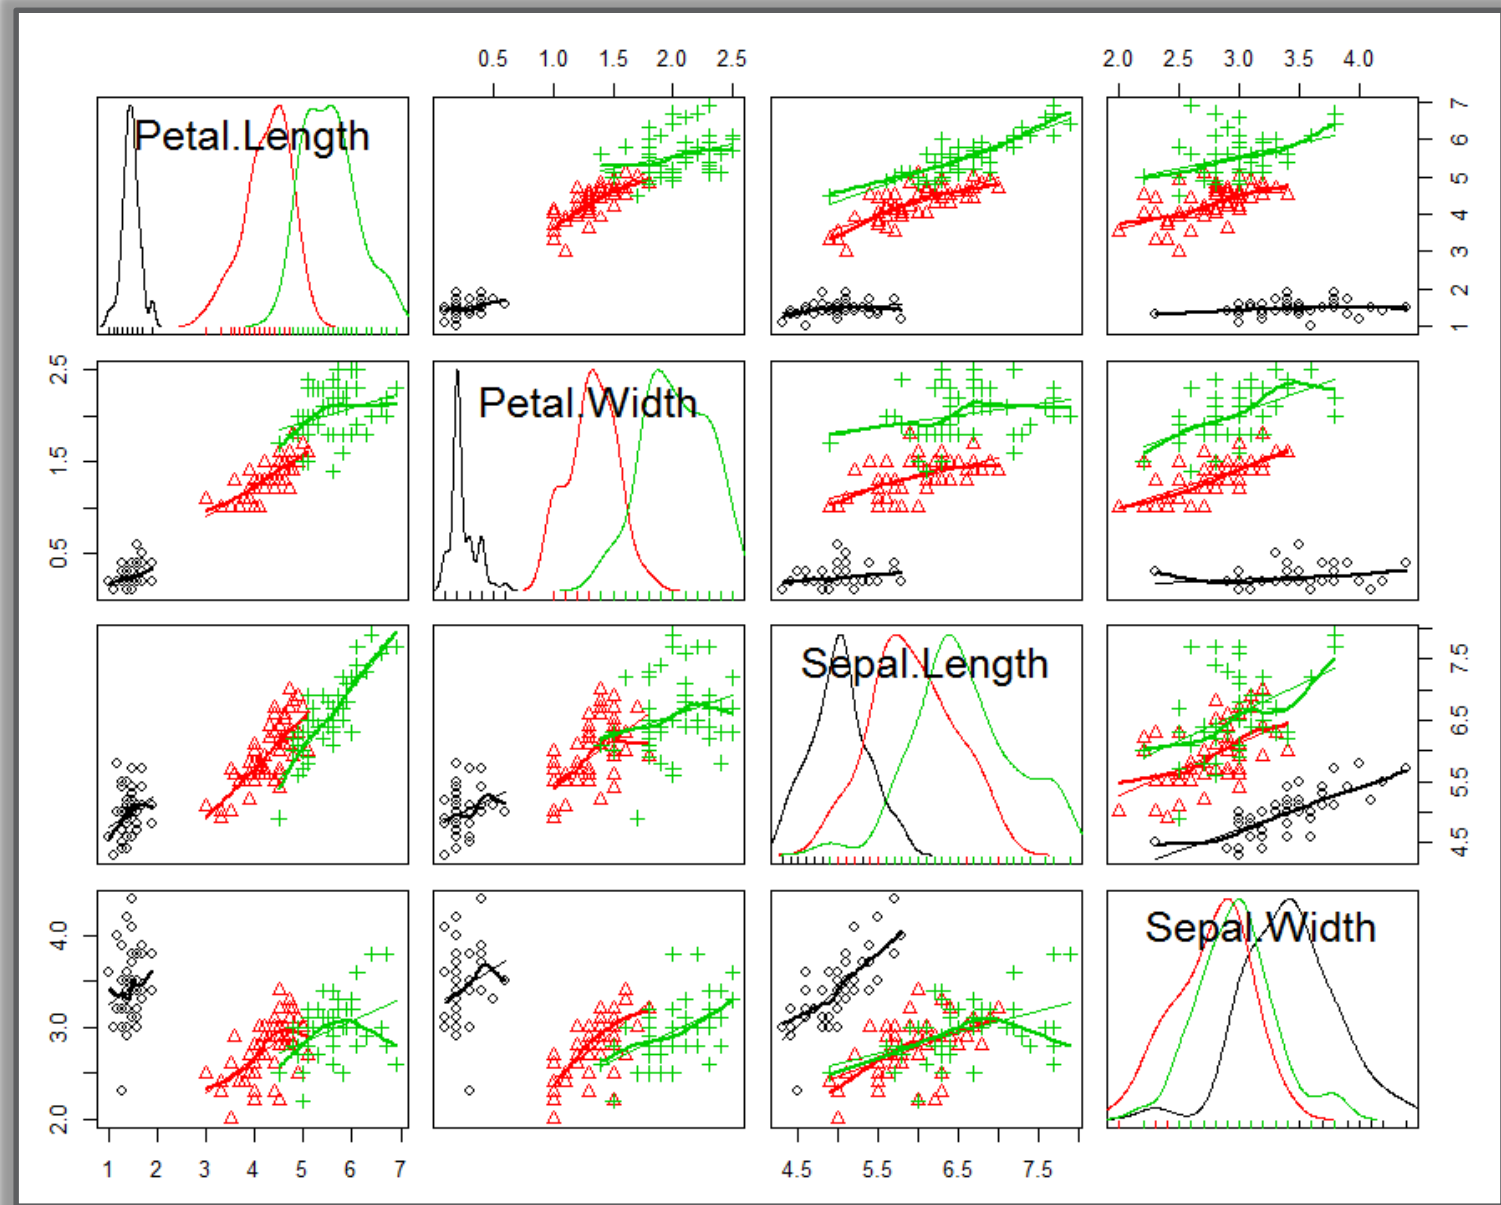
Importance of components:

	Comp.1	Comp.2	Comp.3	Comp.4
Standard deviation	1.7083611	0.9560494	0.38308860	0.143926497
Proportion of Variance	0.7296245	0.2285076	0.03668922	0.005178709
Cumulative Proportion	0.7296245	0.9581321	0.99482129	1.000000000




Original Relationships: Features are **Correlated!**

Scatterplot Matrix



Extracted Principal Components (PCs) are **Uncorrelated**

	PC1	PC2	PC3	PC4
PC1	1e+00	-1e-16	-1e-15	2e-15
PC2	-1e-16	1e+00	2e-16	-5e-16
PC3	-1e-15	2e-16	1e+00	-1e-15
PC4	2e-15	-5e-16	-1e-15	1e+00



The new extracted features (PCs) are **UNCORRELATED!**

PC is a weighted linear sum of original features

Extracted Feature → PCA-based Feature Extraction:

$$PC = w_1 * f_1 + w_2 * f_2 + \dots + w_d * f_d$$

The magnitude of each **weight** indicates **how important the corresponding feature is** → it could be used as a **feature selection** technique!

	PC1	PC2	PC3	PC4
Sepal.Length	0.36	-0.66	0.58	0.3
Sepal.Width	-0.08	-0.73	-0.60	-0.3
Petal.Length	0.86	0.17	-0.08	-0.5
Petal.Width	0.36	0.08	-0.55	0.8

$$PC1 = 0.86 \text{ Petal.Length} + 0.36 \text{ Petal.Width} \\ + 0.36 \text{ Sepal.Length} - 0.08 \text{ Sepal.Width}$$

$$PC2 = -0.73 \text{ Sepal.Width} - 0.66 \text{ Sepal.Length} \\ + 0.17 \text{ Petal.Length} + 0.08 \text{ Petal.Width}$$

Preserved Variability for **Top-k PCs**: $d \rightarrow k$ ($k \ll d$)

Percentage of variability preserved
if the first k PCs are used for projection:

$$\frac{\sum_{i=1}^k \lambda_i}{\sum_{j=1}^d \lambda_j}$$

Standard deviation = sqrt (Variance)


```
Standard deviations (1, ..., p=4):  
[1] 2.06 0.49 0.28 0.15
```

Variance preserved by PC1:

$$2.06 \times 2.06 = 4.2$$

SORTED Eigenvalues: `pca.model$sdev`

$$\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_k \geq \dots \geq \lambda_d \geq 0$$



Eigenvalues: Importance of Eigenvectors

SORTED:

$$\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_k \geq \dots \geq \lambda_d \geq 0$$

Importance of Principal Components (PC), or Eigenvectors, or Rotations

PC1 preserves more variance than PC2

PC2 preserves more variance than PC3

....


Proportion of Variance Preserved if only k PCs are used:

$$\frac{\sum_{i=1}^k \lambda_i}{\sum_{j=1}^d \lambda_j}$$

Principal Components (eigenvectors) are NOT unique

PCA: **Linear Orthogonal** Transformation:

- PCs are simply a rotation of the original coordinate system
- Each PC is a **line**
- Each PC is **perpendicular** to the other PCs
 - PCs are **Uncorrelated**
 - The angle between any pair of PCs is 90°




PCA: Linear Orthogonal Dimension Reduction

Principal Component Analysis (**PCA**) finds **intrinsic** dimensionality and allows for low-dimensional representation of the data.



Principal Component Analysis is the **Spectral Decomposition of the Covariance Matrix**.



PCA: Covariance vs. Correlation Matrix

- PCA results **depend on the scales** at which the variables are measured:
 - PCA should only be used with the raw data if all variables have the same units of measure
- PCA results **depend on the variances** of the variables: the ones with the highest sample variances will tend to be emphasized in the first few principal components:
 - Use PCA with covariance matrix only if you wish to give variables with higher variances more weight in the analysis
- If the variables either have **different units of measurement** (i.e., pounds, feet, gallons, etc), or if we wish **each variable to receive equal weight in the analysis**, then the variables should be **standardized (Z-scores)** before a principal components analysis is carried out.