

Introduction to Machine Learning

Ranga Raju Vatsavai, Ph.D.

Chancellors Faculty Excellence Associate Professor in Geospatial Analytics
Department of Computer Science, North Carolina State University (NCSU)

Feb. 25-27, 2019

Today

- Nearest Neighbor Classification


2/7/16

Types of Learners

- Eager learner
 - Designed to learn a model from instances (examples) that maps input attributes to the class label as soon as training data is available
 - Decision trees, rule-based, SVM, ...
- Lazy learner
 - Opposite of eager learner. Delay the process of learning a model from training data until it is needed to classify a new sample
 - Rote-learner
 - Memorizes entire training data and performs classification only if attributes of record match one of the training examples exactly
 - Nearest neighbor
 - Uses k “closest” instances (nearest neighbors) for performing classification

Nearest Neighbor Classifiers


- Basic idea:
 - If it walks like a duck, quacks like a duck, then it’s probably a duck



Basic k-nearest neighbor algorithm


- Training method
 - Save the training examples
- Prediction
 - Find the k training examples $(x_1, y_1), \dots, (x_k, y_k)$ that are closest to the test example $(x, ?)$
 - Assign the most frequent class among those y_i 's

Nearest Neighbor Classifiers



- Requires three things
 - The set of labeled records
 - Distance Metric to compute distance between records
 - The value of k , the number of nearest neighbors to retrieve
- To classify an unknown record:
 - Compute distance to other training records
 - Identify k nearest neighbors
 - Use class labels of nearest neighbors to determine the class label of unknown record (e.g., by taking majority vote)

Definition of Nearest Neighbor




(a) 1-nearest neighbor (b) 2-nearest neighbor (c) 3-nearest neighbor

K -nearest neighbors of a record x are data points that have the k smallest distances to x

1 nearest neighbor

- Voronoi Diagram



Nearest Neighbor Classification

- Compute distance between two points:
 - Euclidean distance

$$d(p, q) = \sqrt{\sum_i (p_i - q_i)^2}$$


- Determine the class from nearest neighbor list
 - Take the majority vote of class labels among the k-nearest neighbors
 - Weigh the vote according to distance
 - weight factor, $w = 1/d^2$

Nearest Neighbor Classification...

- Choosing the value of k:

- If k is too small, sensitive to noise points

- If k is too large, neighborhood may include points from other classes



Nearest Neighbor Classification...

- Scaling issues
 - Attributes may have to be scaled to prevent distance measures from being dominated by one of the attributes
 - Example:
 - height of a person may vary from 1.5m to 1.8m
 - weight of a person may vary from 90lb to 300lb
 - income of a person may vary from \$10K to \$1M

Nearest Neighbor Classification...

- Selection of the right similarity measure is critical:

1 1 1 1 1 1 1 1 1 0	vs	0 0 0 0 0 0 0 0 0 1
0 1 1 1 1 1 1 1 1 1		1 0 0 0 0 0 0 0 0 0

Euclidean distance = 1.4142 for both pairs

Distance Measure

- Euclidian
 - Why its not best?
 - Distances in each dimension are squared before summation places great emphasis on those features which dissimilarity is large
 - How about just absolute differences
 - Manhattan or city block or taxi-cab distance
 - How would you deemphasize single large feature difference and more influenced by numerous small ones?

Nearest neighbor Classification...

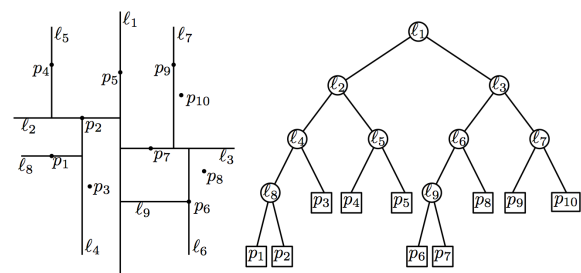
- k-NN classifiers are lazy learners since they do not build models explicitly
- Classifying unknown records are relatively expensive
- Can produce arbitrarily shaped decision boundaries
- Easy to handle variable interactions since the decisions are based on local information
- Selection of right proximity measure is essential
- Superfluous or redundant attributes can create problems
- Missing attributes are hard to handle

Improving KNN Efficiency

- Avoid having to compute distance to all objects in the training set
 - Multi-dimensional access methods (k-d trees)
 - Fast approximate similarity search
 - Locality Sensitive Hashing (LSH)
- Condensing
 - Determine a smaller set of objects that give the same performance
- Editing
 - Remove objects to improve efficiency

Improving KNN Efficiency

- Multi-dimensional access methods (k-d trees)



Acknowledgements: UFL Database Group

Improving KNN Efficiency

- How is it going to improve performance?

Acknowledgements: UFL Database Group

Practice

- Compare KNN with K-Means
 - Assume that the data is labeled, so we can label clusters
 - Given a new point, how each method predicts